


Module 3

NoSQL Databases

María del Mar Roldán – University of Málaga
Ismael Navas – University of Málaga



Learning Objectives

- Understand the Cassandra data model
- Introduce cqlsh
- Understand and use the DDL subset of CQL
- Introduce DevCenter
- Understand and use the DML subset of CQL
- **Understand basics of data modeling**

What is data modeling?

- Data modeling is a process that involves
 - Collection and analysis of data requirements in an information system
 - Identification of participating entities and relationships among them
 - Identification of data access patterns
 - A particular way of organizing and structuring data
 - Design and specification of a database schema
 - Schema optimization and data indexing techniques
- Data modeling = Science + Art

Steps for data modeling in Cassandra

1. Understand data and application queries

- Data may or may not exist in some format (RDBMS, XML, CSV, ...)
- Queries can be organized into a query graph

2. Design column families

- Design is based on access patterns or queries over data

3. Implement the design using CQL

- Optimizations concerning data types, keys, partition sizes, ordering

Key steps for data modeling

The products of the data modeling steps are documented as

- Conceptual data model
 - Technology-independent, unified view of data
 - Entity-relationship model, dimensional model, etc.
- Logical data model
 - Unique for Cassandra
 - Column family diagrams
- Physical data model
 - Unique for Cassandra
 - CQL definitions

Apache Cassandra

Is relational database design similar to
Cassandra database design?

DATASTAX

No!

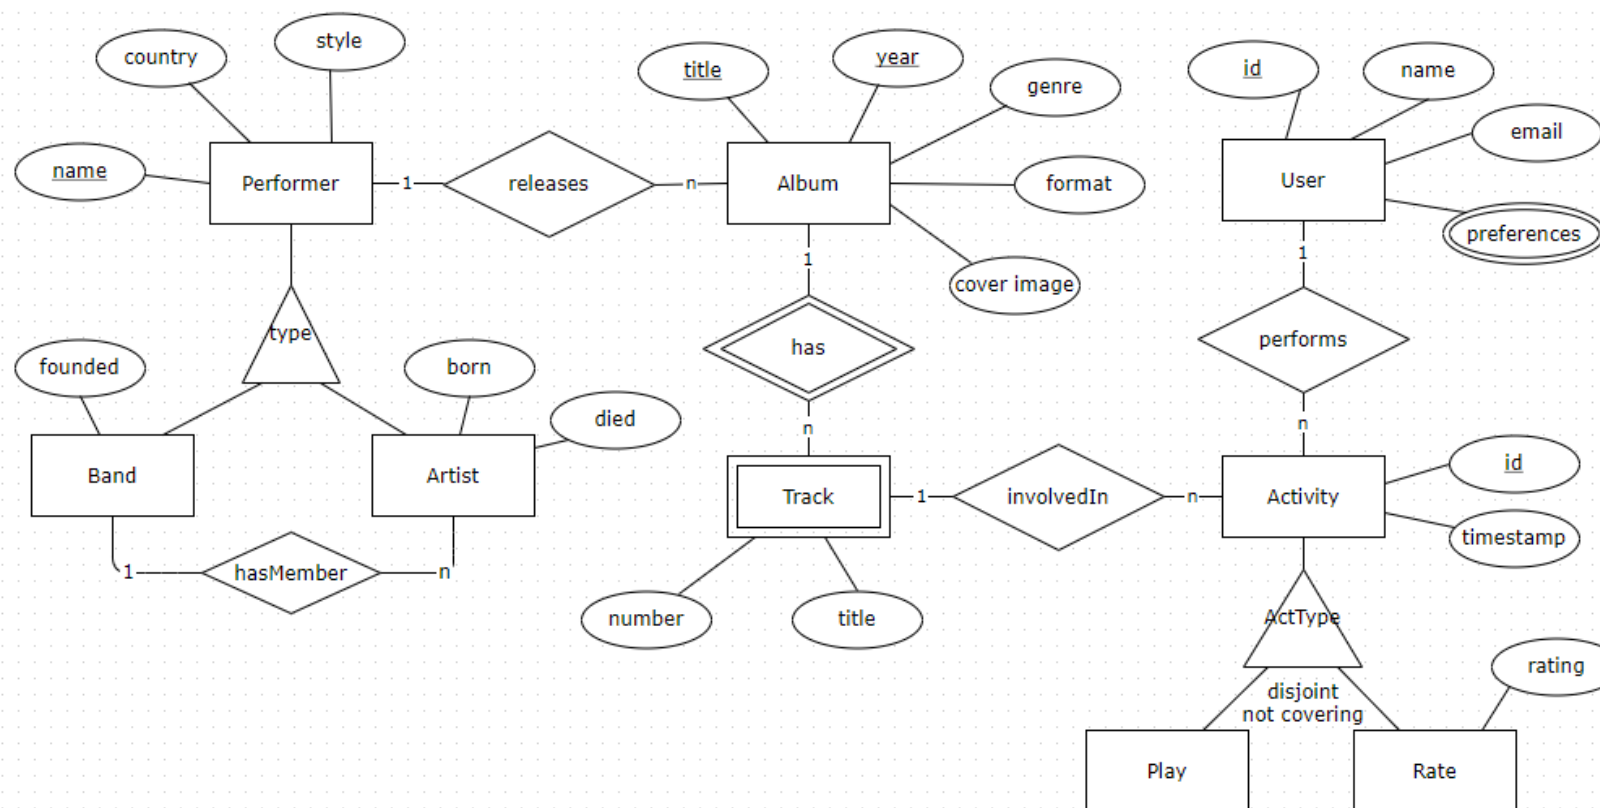
Cassandra

- Multi-dimensional column family
 - Equally good for simple and complex data
- All data required to answer a query must be nested in a column family
 - Referential integrity is a non-issue
- Data modeling methodology is driven by queries and data
 - Data duplication is considered normal (side effect of data nesting)

Relational

- Two-dimensional relation
 - Suited for simple data
 - Complex data requires many relations and “star” schemas
- Data from many relations is combined to answer a query
 - Referential integrity is important
- Data modeling is driven by data only
 - Data duplication is considered a problem (normalization theory)

Example: Modeling musicdb Database




Example: musicdb Queries

- Q1: find information for a specified user.
- Q2: find performers for a specified style; order by performer (ASC).
- Q3: find information for a specified performer (artist or band)
- Q4: find information (performer, genre, tracks title) for a specified album (given its title and year)
- Q5: find albums (year, title and genre) for a specified performer; order by album release year (DESC) and title (ASC)
- Q6: find albums and performers for a specified genre; order by performer (ASC), year (DESC), and title (ASC).
- Q7: find activities for a specified user; order by activity time (DESC).
- Q8: find statistics for a specified track
- Q9: find user activities for a specified activity type

Example: musicdb Access Patterns

- Q1: find information for a specified user.


GIVEN	FIND
USER.ID	USER.NAME
	USER.EMAIL
	USER.PREFERENCES



Example: musicdb Logical Model

- Q1: find information for a specified user.

GIVEN	FIND
USER.ID	USER.NAME
	USER.EMAIL
	USER.PREFERENCES



USERS_BY_ID

USER_ID NUMBER

USER_NAME TEXT

EMAIL TEXT


PREFERENCES SET<TEXT>

Example: musicdb Physical Model

- Q1: find information for a specified user.

USERS_BY_ID

ID NUMBER
NAME TEXT
EMAIL TEXT
PREFERENCES SET<TEXT>



USERS_BY_ID

USER_ID UUID
USER_NAME VARCHAR
EMAIL VARCHAR
PREFERENCES SET<VARCHAR>
PRIMARY KEY(USER_ID)

SELECT * FROM USERS_BY_ID WHERE USER_ID = ?

Example: musicdb CQL script


- Q1: find information for a specified user.

```
CREATE TABLE users_by_id (  
    user_id UUID,  
    user_name VARCHAR,  
    email VARCHAR,  
    preferences SET<VARCHAR>,  
    PRIMARY KEY (user_id)  
);
```

Example: musicdb Access Patterns

- Q2: find performers for a specified style; order by performer (ASC).

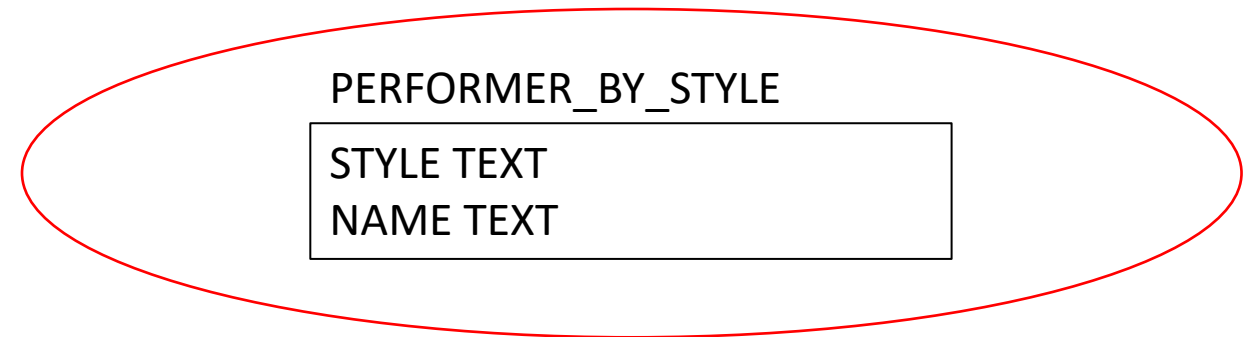

GIVEN	FIND
PERFORMER.STYLE	NAME (ASC)



Example: musicdb Logical Model

- Q2: find performers for a specified style; order by performer (ASC).

GIVEN	FIND
PERFORMER.STYLE	PERFORMER.NAME (ASC)




Example: musicdb Physical Model

- Q2: find performers for a specified style; order by performer (ASC).

PERFORMER_BY_STYLE

STYLE TEXT
NAME TEXT



PERFORMER_BY_STYLE

STYLE VARCHAR
NAME VARCHAR
PRIMARY KEY (STYLE, NAME)

Partition key

Clustering column

SELECT NAME FROM PERFORMER_BY_STYLE WHERE STYLE = ?

Example: musicdb CQL script

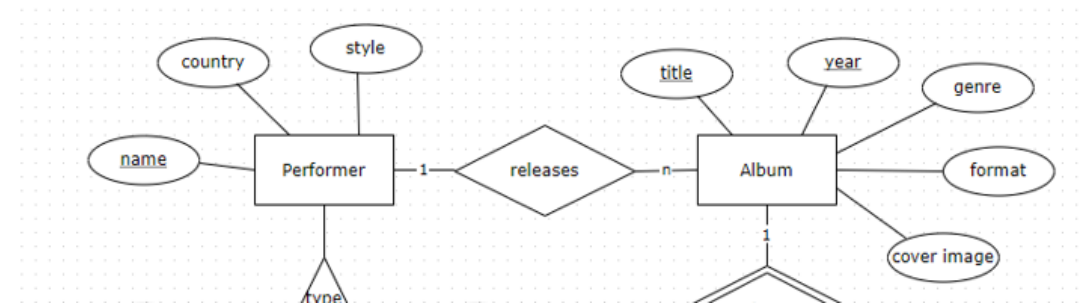
- Q2: find performers for a specified style; order by performer (ASC).

```
CREATE TABLE performer_by_style(  
    style VARCHAR,  
    name VARCHAR,  
    PRIMARY KEY (style, name)  
);
```


Example: musicdb Access Patterns

- Q6: find albums and performers for a specified genre; order by performer (ASC), year (DESC), and title (ASC).

GIVEN	FIND
ALBUM.GENRE	PERFORMER.NAME (ASC)
	ALBUM.YEAR (DESC)
	ALBUM.TITLE (ASC)



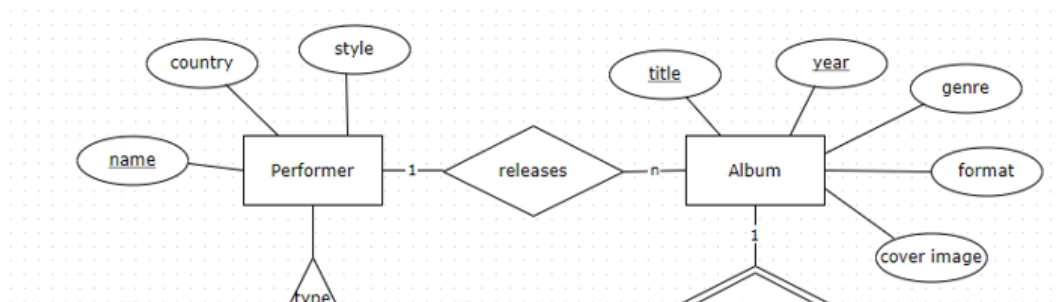
Example: musicdb Logical Model

- Q6: find albums and performers for a specified genre; order by performer (ASC), year (DESC), and title (ASC).

GIVEN	FIND
ALBUM.GENRE	PERFORMER.NAME (ASC)
	ALBUM.YEAR (DESC)
	ALBUM.TITLE (ASC)

ALBUMS_BY_GENRE

GENRE TEXT
PERFORMER TEXT
YEAR INT
TITLE TEXT




Example: musicdb Physical Model

- Q6: find albums and performers for a specified genre; order by performer (ASC), year (DESC), and title (ASC).

ALBUMS_BY_GENRE

GENRE TEXT
PERFORMER TEXT
YEAR INT
TITLE TEXT



ALBUMS_BY_GENRE

GENRE VARCHAR
PERFORMER VARCHAR
YEAR INT
TITLE VARCHAR
PRIMARY KEY(GENRE, PERFORMER, YEAR, TITLE)

Partition key


Clustering columns

SELECT PERFORMER, YEAR, TITLE FROM ALBUMS_BY_GENRE WHERE GENRE = ?

Example: musicdb CQL script

- Q6: find albums and performers for a specified genre; order by performer (ASC), year (DESC), and title (ASC).

```
CREATE TABLE albums_by_genre(  
    genre VARCHAR,  
    performer VARCHAR,  
    year INT,  
    title VARCHAR,  
    PRIMARY KEY (genre, performer, year, title)  
) WITH CLUSTERING ORDER BY (performer ASC, year DESC, title ASC);
```

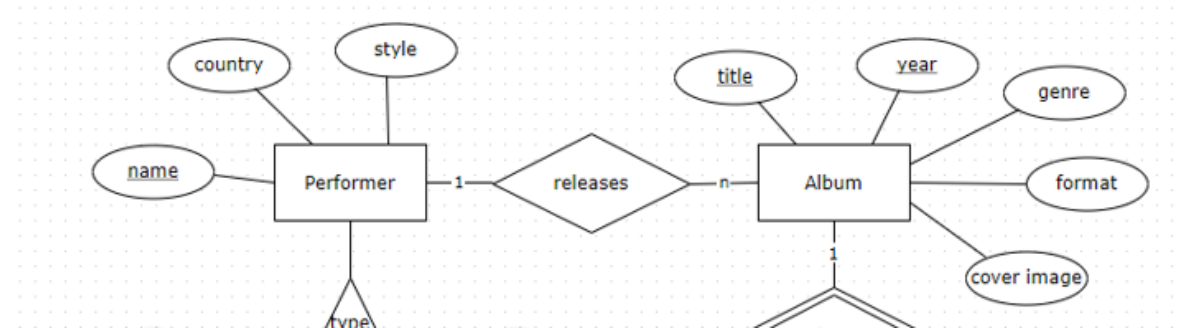


Time for Exercises

Example: musicdb Access Patterns

- Q5: find albums (year, title and genre) for a specified performer; order by album release year (DESC) and title (ASC)

GIVEN	FIND
Performer.name	Album.genre
	Album.year
	Album.title




Example: musicdb Access Patterns

- Q5: find albums (year, title and genre) for a specified performer; order by album release year (DESC) and title (ASC)

GIVEN	FIND
Performer.name	Album.genre
	Album.year
	Album.title


ALBUMS_BY_PERFORMER

Performer text
Genre text
Year number
Title text



Example: musicdb Access Patterns

- Q5: find albums (year, title and genre) for a specified performer; order by album release year (DESC) and title (ASC)



ALBUMS_BY_PERFORMER

Perfomer text
Genre text
Year number
Title text

ALBUMS_BY_PERFORMER

performer varchar
genre varchar
year int
title varchar
Primary key(performer, year, title)

Example: musicdb Access Patterns

- Q7: find activities for a specified user; order by activity time (DESC).

GIVEN	FIND
User.id	Track.number
	Track. title
	Album.title
	Album.year
	Rate.rating
	Activity.timestamp

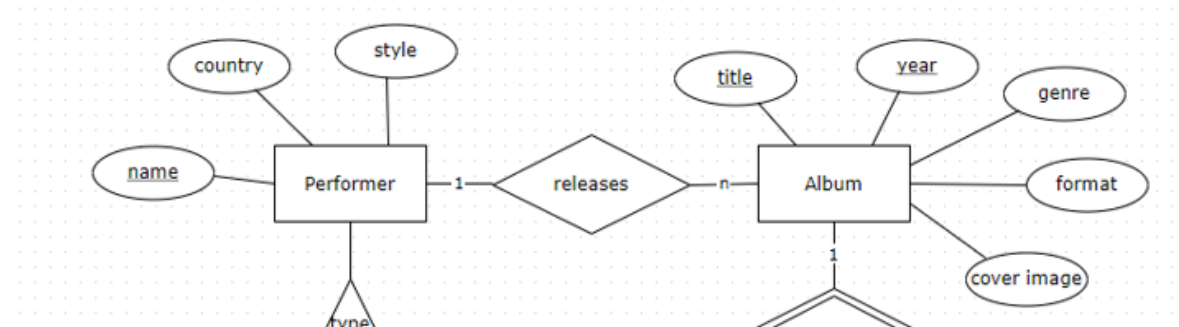
ACTIVITIES_BY_USER

```
User UUID
Track_number int
Album_title varchar
album_year int
Rating int
Activity TIMEUUID
PRIMARY KEY(user, activity)
```

Example: musicdb Access Patterns

- Q5: find albums (year, title and genre) by each year; order by album title (ASC) and performer (ASC)

GIVEN	FIND




Example: musicdb Access Patterns

- Q5: find albums (year, title and genre) by each year; order by album title (ASC) and performer (ASC)

GIVEN	FIND
	Album.title
	Album.year
	Album.genre

↑
EMPTY GIVEN!!!!




Example: musicdb Access Patterns

- Q5: find albums (year, title and genre) by each year; order by album title (ASC) and performer (ASC)

GIVEN	FIND
	Album.year
	Album.title
	Album.genre


ALBUMS_PER_YEAR

year number
title text
genre text



Example: musicdb Access Patterns

- Q5: find albums (year, title and genre) by each year; order by album title (ASC) and performer (ASC)



ALBUMS_PER_YEAR

```
year int
title varchar
genre varchar
Primary key(year, title, performer)
```

ALBUMS_PER_YEAR

```
year number
title text
genre text
```

```
SELECT * FROM ALBUMS_PER_YEAR;
```

```
SELECT * FROM ALBUMS_PER_YEAR where title < ... allow filtering;
```

```
SELECT * FROM ALBUMS_PER_YEAR where title = ... and genre = ... allow filtering;
```