



Máster de Formación Permanente  
en **BIG DATA**  
e Inteligencia Artificial

Máster en

**Big Data, Inteligencia Artificial e Ingeniería de Datos**

---



# Module 4

## NoSQL Databases

María del Mar Roldán – University of Málaga

## Table of contents

- Cassandra



Máster de Formación Permanente  
en **BIG DATA**  
e Inteligencia Artificial

Máster en

**Big Data, Inteligencia Artificial e Ingeniería de Datos**

---



# Cassandra

## Definition of Cassandra

Apache Cassandra™ is a free

Distributed...

High performance...

Extremely scalable...

Fault tolerant (i.e. no single point of failure)...



post-relational database solution. Cassandra can serve as both real-time datastore (the “system of record”) for online/transactional applications, and as a read-intensive database for business intelligence systems.


# Apache Cassandra- Introducción

- Web: <http://cassandra.apache.org/>
- Documentación: <http://www.datastax.com/docs/>
- Cassandra use cases
  - Digg
  - Netflix
  - Rackspace
  - Twitter
  - ...




# Architecture Overview

- Cassandra was designed with the understanding that system/hardware failures can and do occur
- Peer-to-peer, distributed system
- All nodes the same
- Data partitioned among all nodes in the cluster
- Custom data replication to ensure fault tolerance
- Read/Write-anywhere design



# Architecture Overview

- Each node communicates with each other through the Gossip protocol, which exchanges information across the cluster every second
- A commit log is used on each node to capture write activity. Data durability is assured
- Data also written to an in-memory structure (memtable) and then to disk once the memory structure is full (an SStable)



# Learning Objectives

- Understand the Cassandra data model
- Introduce cqlsh
- Understand and use the DDL subset of CQL
- Introduce DataGrip
- Understand and use the DML subset of CQL
- ~~Understand basics of data modeling (Challenge)~~



## Learning Objectives

- **Understand the Cassandra data model**
- Introduce cqlsh
- Understand and use the DDL subset of CQL
- Introduce DataGrip
- Understand and use the DML subset of CQL
- ~~Understand basics of data modeling (Challenge)~~

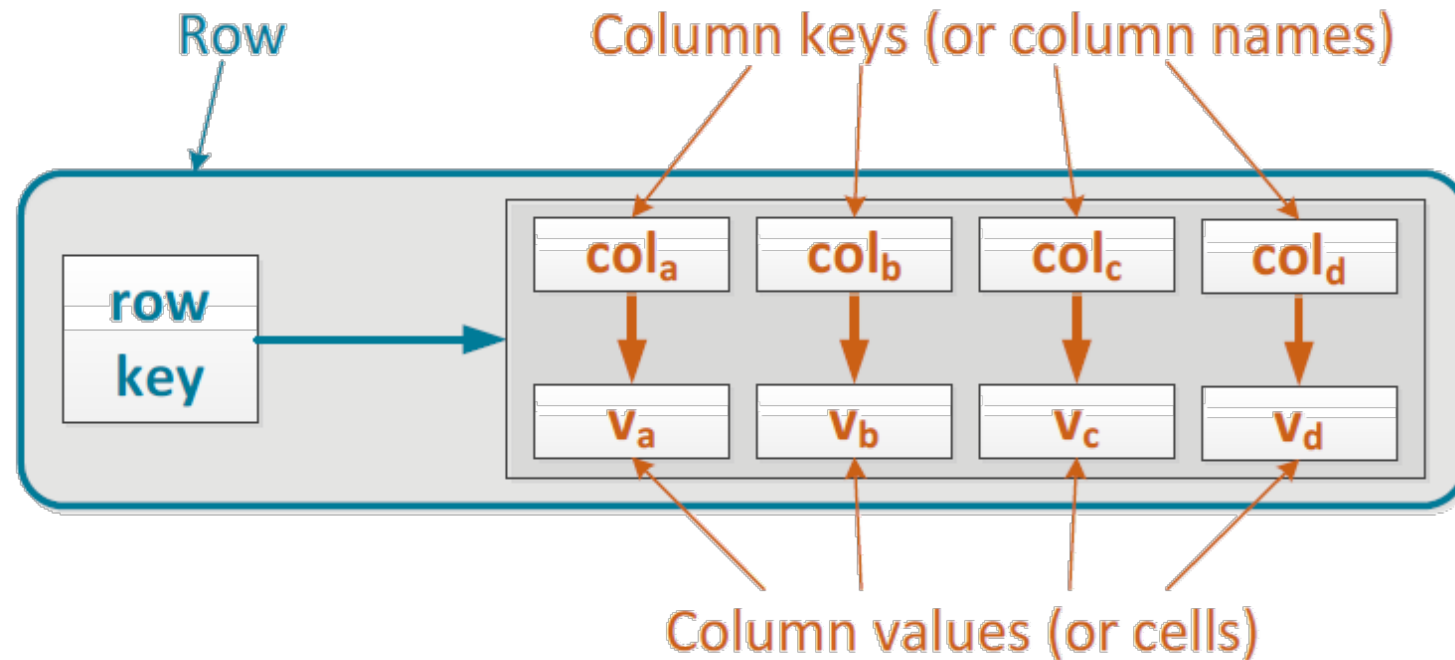
# What are the essential constituents of the Cassandra data model?

- The Cassandra data model defines
  1. *Column family* as a way to store and organize data
  2. *Table* as a two-dimensional view of a multi-dimensional *column family*
  3. Operations on tables using the Cassandra Query Language (CQL)
- We cover these three constituents in the order they are listed
  - Understanding *column families* is a prerequisite to understanding *tables*
  - Understanding *tables* is a prerequisite to understanding operations

# What are row, row key, column key, and column value?

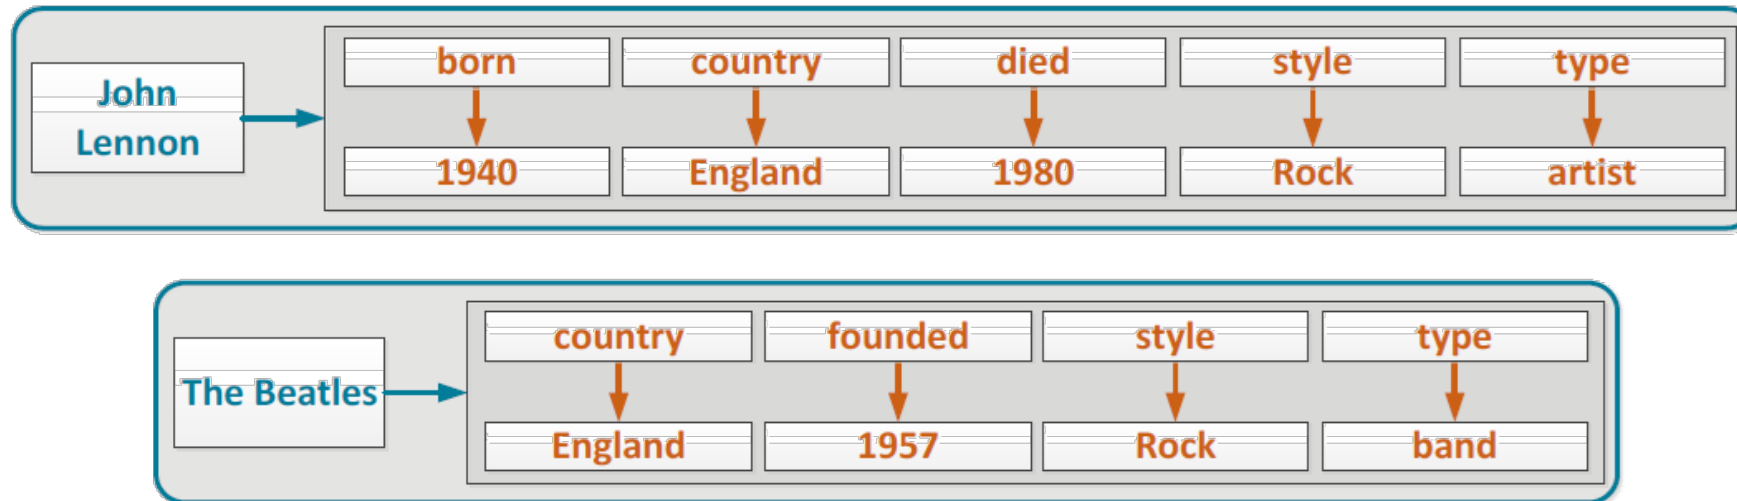
- *Row* is the smallest unit that stores related data in Cassandra
  - Rows – individual rows constitute a *column family*
  - Row key – uniquely identifies a *row* in a *column family*
  - Row – stores pairs of *column keys* and *column values*
  - Column key – uniquely identifies a *column value* in a *row*
  - Column value – stores one value or a *collection* of values

# Apache Cassandra – Fila (Row)




# What are row, row key, column key, and column value?

- Sample rows that describe an artist and a band
  - *Column keys* are inherently sorted
  - A *row* can be retrieved if its *row key* is known
  - A *column value* can be retrieved if its *row key* and *column key* are known



# Apache Cassandra – Fila (Row)

- Rows may be described as “skinny” or “wide”
  - Skinny row – has a fixed, relatively small number of *column keys*
    - Previous examples were skinny rows
  - Wide row – has a relatively large number of *column keys* (hundreds or thousands); this number may increase as new data values are inserted
    - For example, a row that stores all bands of the same style
    - The number of such bands will increase as new bands are formed




- Note that column values do not exist in this example
  - The column key – in this case a band name – stores all the data desired
  - Could have stored the number of albums, or year founded, etc., as column values

# What are composite row key and composite column key?


- Composite row key – multiple components separated by colon

- 'Revolver' and 1966 are the album title and year
- 'tracks' value is a collection (map)




- Composite column key – multiple components separated by colon

- Composite column keys are sorted by each component
- 1,2, ..., 14 are track numbers; 'title' is metadata
  - We could have stored actual title as components of composite column keys: 1:Taxman, 2:Eleanor Rigby, ..., 14:Tomorrow Never Knows



# Can simple and composite column keys co-exist in the same row?

- Row can contain both simple and composite column keys
  - 'genre' and 'performer' are simple column keys
  - '1:title', '2:title', ... are composite column keys



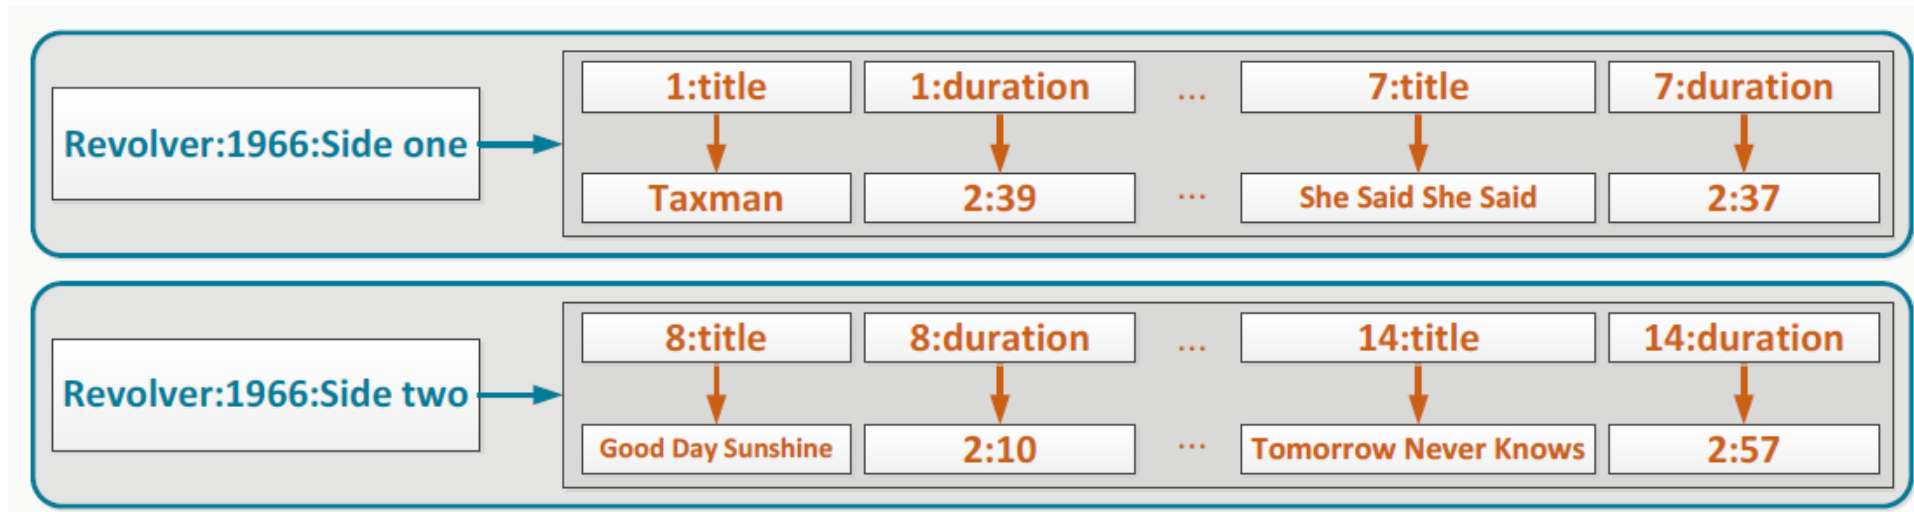


# What components of a row can store useful values?

- Any component of a row can store data or metadata
  - Simple or composite row keys
  - Simple or composite column keys
  - Atomic or set-valued (collection) column values
    - Metadata: 'Side one', 'Side two', 'title', 'duration'
    - Data: everything else ('Revolver', '1966', 'She Said She Said', etc.)

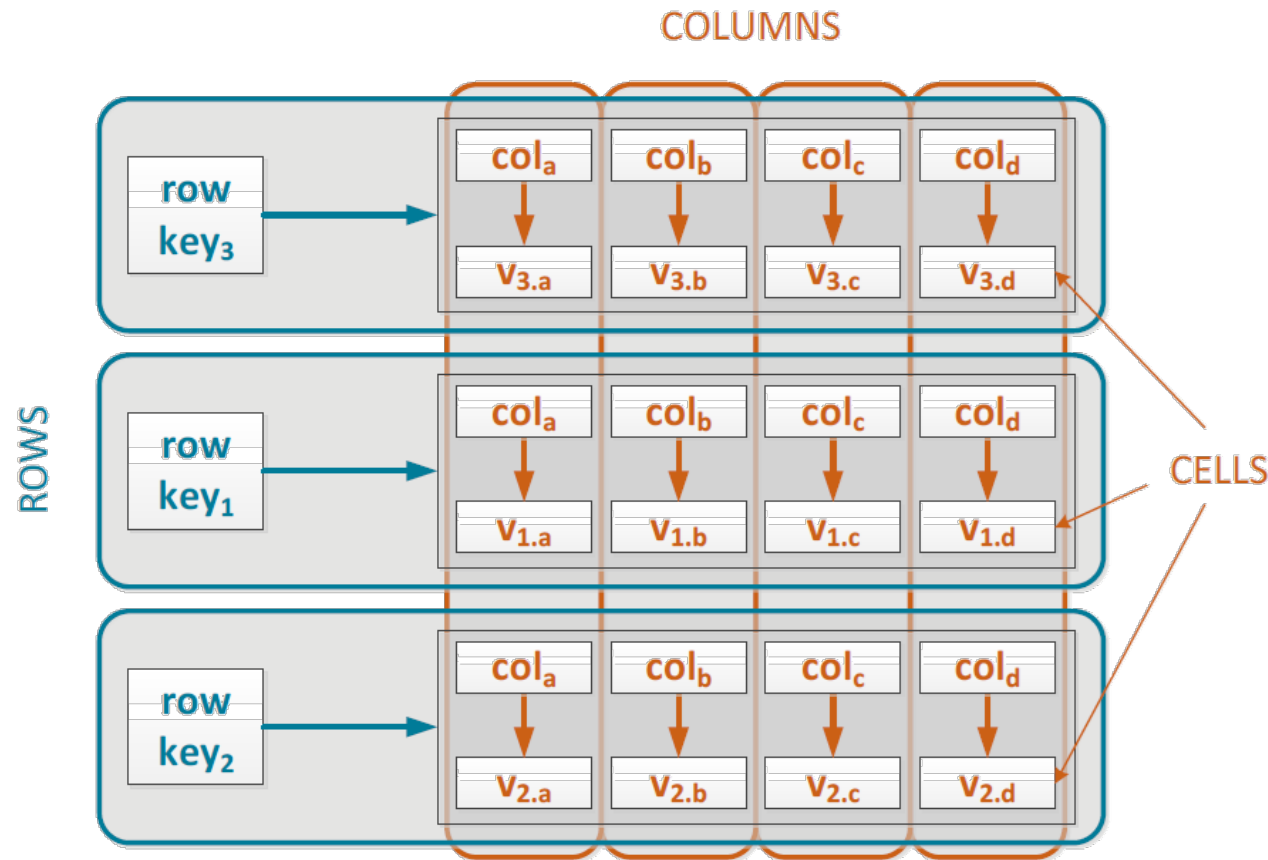
# What components of a row can store useful values?

- Metadata: 'Side one', 'Side two', 'title', 'duration'
- Data: everything else ('Revolver', '1966', 'She Said She Said', etc.)



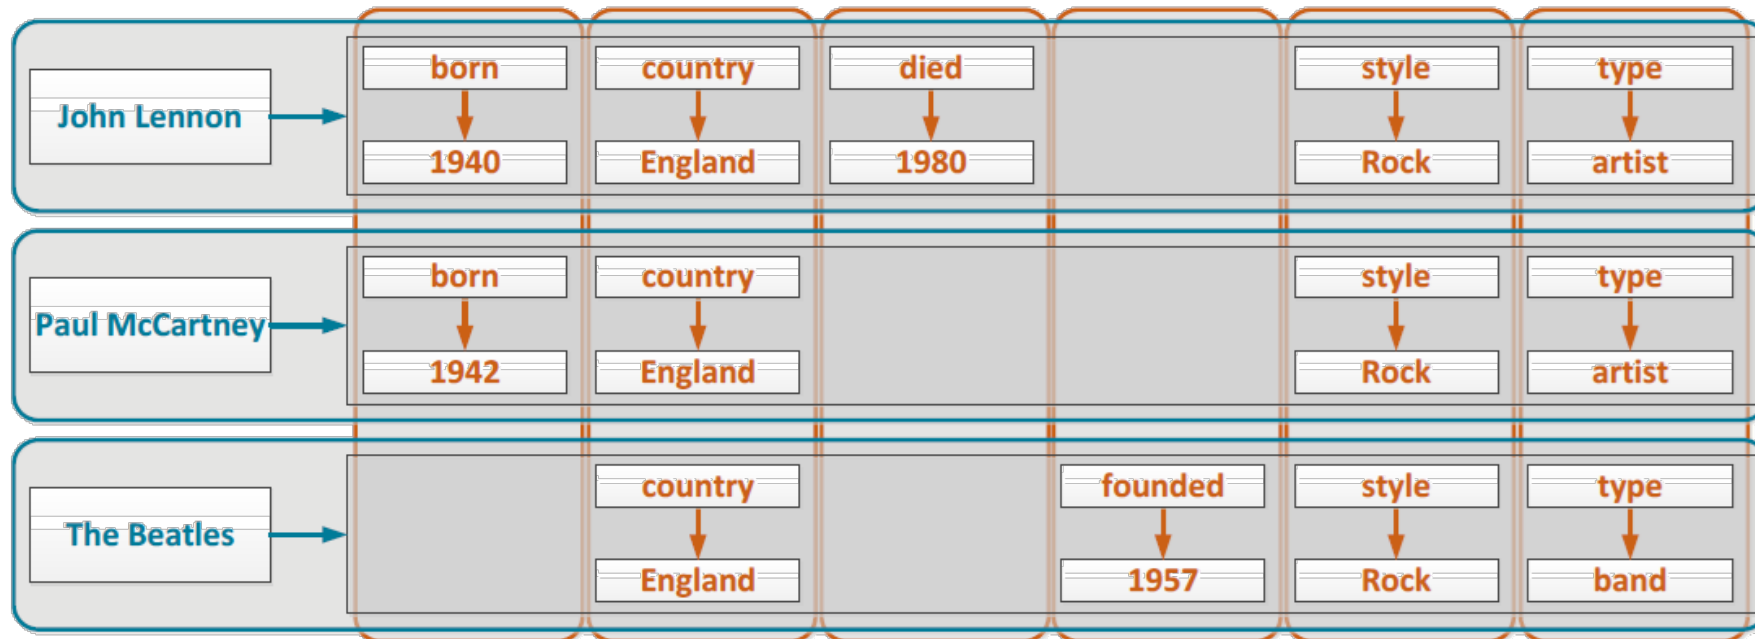
# Apache Cassandra – Column Family

- Column family – set of rows with a similar structure



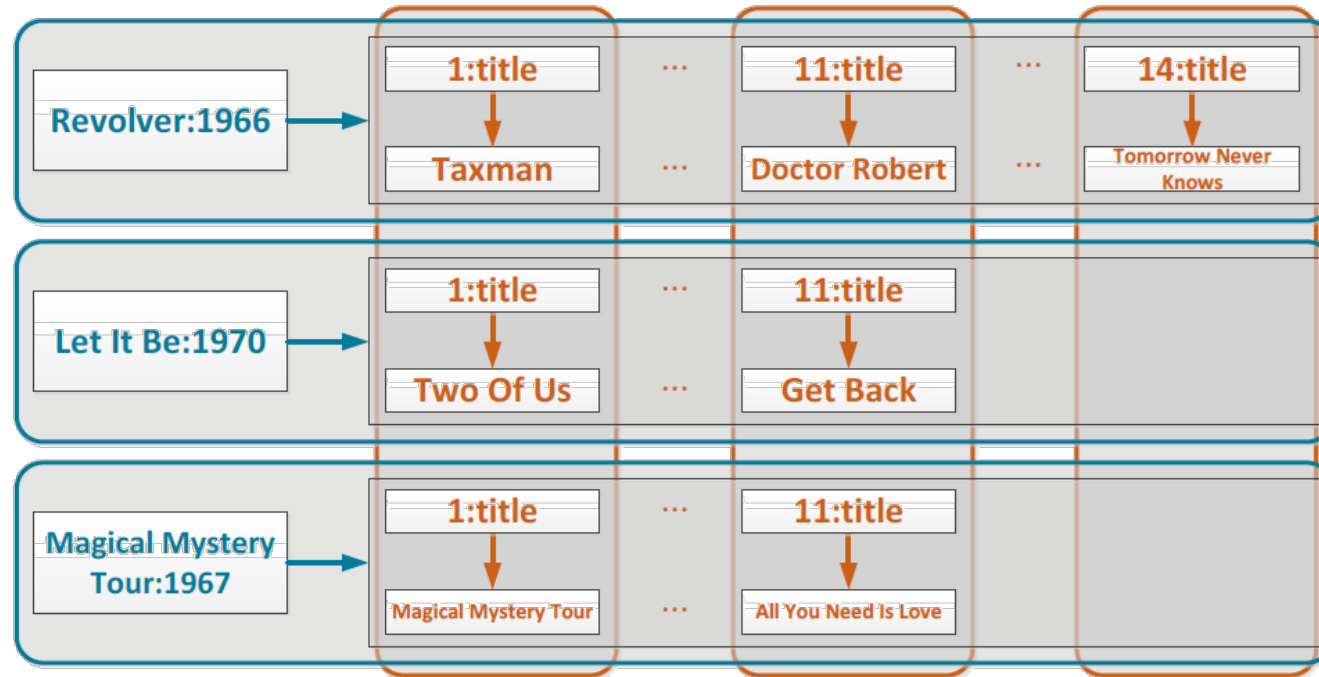
# Apache Cassandra – Column Family

- Distributed
- Sparse
  - Column family that stores data about artists and bands



# Apache Cassandra – Column Family

- Sorted columns
- Multidimensional
  - Column family that stores albums and their tracks



## Apache Cassandra – Column Family

- Size of a column family is only limited to the size of a cluster
  - Linear scalability
  - **Rows are distributed among the nodes in a cluster**
- Column family component size considerations
  - Data from a one row must fit on one node
    - Data from any given row never spans multiple nodes
  - Maximum columns per row is 2 billion
    - In practice – Up to 100 thousand
  - Maximum data size per cell (column value) is 2 GB
    - In practice – Up to 100 MB

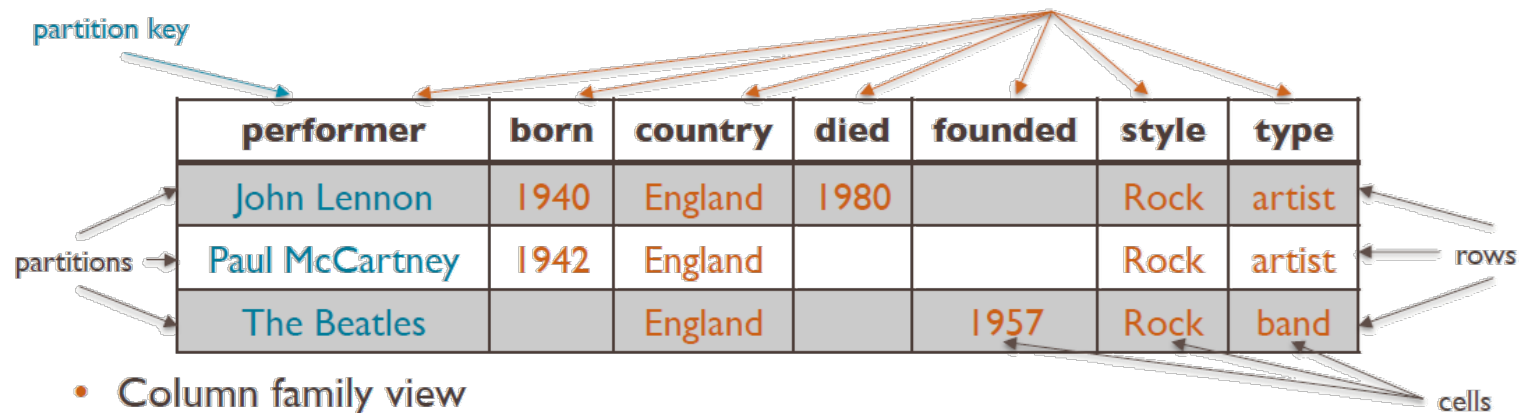
## Apache Cassandra – CQL Table

- What is a CQL table and how is it related to a column family?
  - **A CQL table is a column family**
    - CQL tables provide two-dimensional views of a column family, which contains potentially multi-dimensional data, due to composite keys and collections
  - CQL table and column family are largely interchangeable terms
    - Not surprising when you recall tables and relations, columns and attributes, rows and tuples in relational databases
  - Supported by declarative language Cassandra Query Language
    - Data Definition Language (DDL), subset of CQL
    - SQL-like syntax, but with somewhat different semantics
    - Convenient for defining and expressing Cassandra database schemas

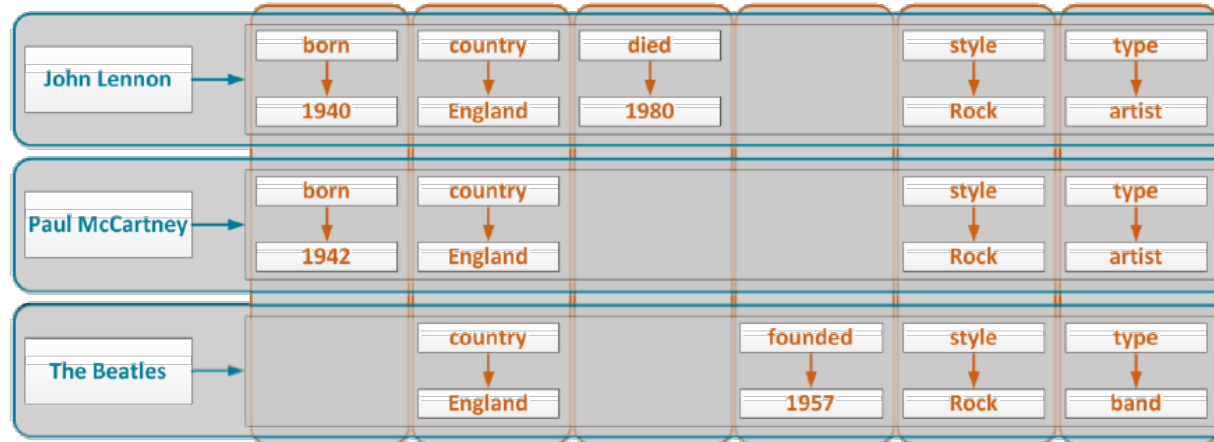


# What are partition, partition key, row, column, and cell?

- Table with single-row partitions



- Column family view

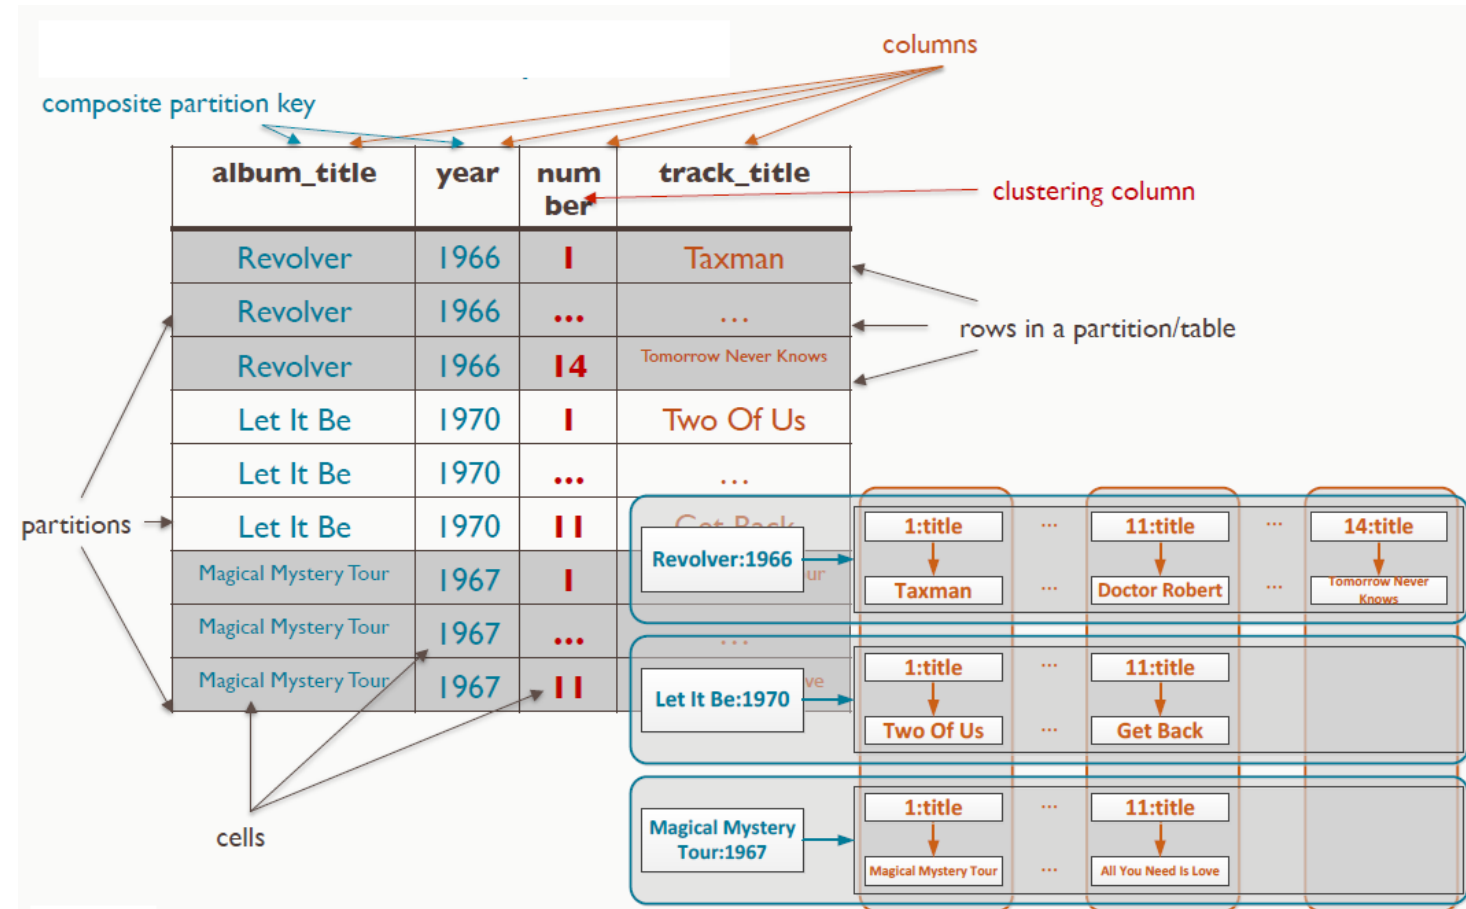






# What are composite partition key and clustering column?

- Table with multi-row partitions



# What are static columns?

- Table with multi-row partitions

composite partition key      clustering column      static columns


album_title	year	number	genre	performer	track_title
Revolver	1966	I	Rock	The Beatles	Taxman
Revolver	1966	...	Rock	The Beatles	...
Revolver	1966	I4	Rock	The Beatles	Tomorrow Never Knows
Let It Be	1970	I	Rock	The Beatles	Two Of Us
Let It Be	1970	...	Rock	The Beatles	...
Let It Be	1970	II	Rock	The Beatles	Get Back
Magical Mystery Tour	1967	I	Rock	The Beatles	Magical Mystery Tour
Magical Mystery Tour	1967	...	Rock	The Beatles	...
Magical Mystery Tour	1967	II	Rock	The Beatles	All You Need Is Love

rows in a partition

cells

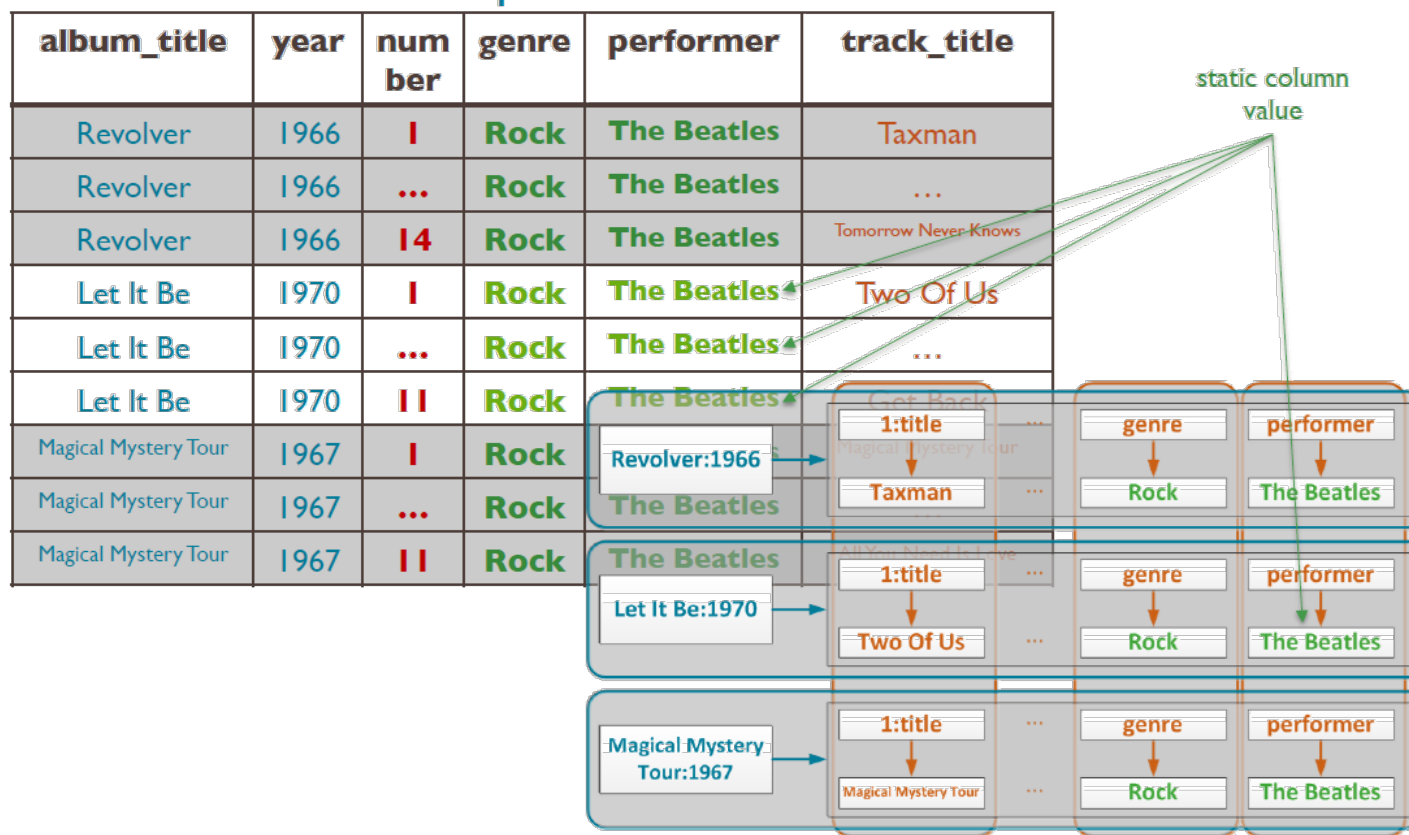
partitions

- Static column values are shared for all rows in a multi-row partition



## What are static columns?

- Table with multi-row partitions



## What is a primary key?

- Primary key uniquely identifies a row in a table
  - Simple or composite partition key and all clustering columns (if present)
  - Primary key
    - performer

performer	born	country	died	founded	style	type
John Lennon	1940	England	1980		Rock	artist
Paul McCartney	1942	England			Rock	artist
The Beatles		England		1957	Rock	band

## What is a primary key?

- Primary key uniquely identifies a row in a table
  - Simple or composite partition key and all clustering columns (if present)
  - Primary key (table above)
    - album, year, number
  - Static columns cannot be part

album_title	year	number	track_title
Revolver	1966	I	Taxman
Revolver	1966	...	...
Revolver	1966	14	Tomorrow Never Knows
Let It Be	1970	I	Two Of Us
Let It Be	1970	...	...
Let It Be	1970	II	Get Back
Magical Mystery Tour	1967	I	Magical Mystery Tour

## What are collection columns?

- Multiple values can be stored in a column
  - Set – typed collection of unique values (e.g., genres)
    - Ordered by values
    - No duplicates
      - {"Blues", "Jazz", "Rock"}
  - List – typed collection of non-unique values (e.g., artists)
    - Ordered by position
    - Duplicates are allowed
      - ["Lennon", "Lennon", "McCartney"]
  - Map – typed collection of key-value pairs (e.g., tracks)
    - Ordered by keys
    - Unique keys but not values
      - {1:"Taxman", 2:"Eleanor Rigby", 3:"I'm Only Sleeping"}

## Apache Cassandra - Map

- Map example
  - Collection column tracks holds a map of album tracks

title	year	genre	performer	tracks
Revolver	1966	Rock	The Beatles	{1: 'Taxman', 2: 'Eleanor Rigby', 3: 'I'm Only Sleeping', 4: 'Love You To', ..., 14: 'Tomorrow Never Knows'}
Let It Be	1970	Rock	The Beatles	{1: 'Two Of Us', 2: 'I Dig A Pony', 3: 'Across The Universe', 4: 'Let It Be', 5: 'Maggie Mae', ..., 11: 'Get Back'}
Magical Mystery Tour	1967	Rock	The Beatles	{1: 'Magical Mystery Tour', 2: 'The Fool On The Hill', 3: 'Flying', 4: 'Blue Jay Way', ..., 11: 'All You Need Is Love'}

## Learning Objectives

- Understand the Cassandra data model
- **Introduce cqlsh**
- Understand and use the DDL subset of CQL
- Introduce DataGrip
- Understand and use the DML subset of CQL
- ~~Understand basics of data modeling (Challenge)~~





## *cqlsh* Shell commands

Command	Description
CAPTURE	Captures command output and appends it to a file
CONSISTENCY	Shows the current consistency level, or given a level, sets it
COPY	Imports and exports CSV (comma-separated values) data
DESCRIBE	Provides information about a Cassandra cluster or data objects
EXPAND	Formats the output of a query vertically
EXIT or QUIT	Terminates cqlsh
SHOW	Shows the Cassandra version, host, or data type assumptions
SOURCE	Executes a file containing CQL statements
TRACING	Enables or disables request tracing

## Learning Objectives

- Understand the Cassandra data model
- Introduce cqlsh
- **Understand and use the DDL subset of CQL**
- Introduce DataGrip
- Understand and use the DML subset of CQL
- ~~Understand basics of data modeling (Challenge)~~

## What is a keyspace or schema?

- Keyspace – a top-level namespace for a CQL table schema
  - Defines the replication strategy for a set of tables
    - Keyspace per application is a good idea
  - Data objects (e.g., tables) belong to a single keyspace
- Replication strategy – the number and pattern by which partitions are copied among nodes in a cluster
  - Two strategies available
    - Simple Strategy (used for prototyping)
    - Network Topology Strategy (production)

# How to create, use and drop keyspaces/schemas?

- To create a keyspace

```
CREATE KEYSPACE musicdb  
WITH replication = {  
  'class': 'SimpleStrategy',  
  'replication_factor' : 3  
};
```

- To assign the working default keyspace for a cqlsh session

```
USE musicdb;
```

- To delete a keyspace and all internal data objects

```
DROP KEYSPACE musicdb;
```

# What is the syntax of the CREATE TABLE statement?

- Primary key declared inline  
CREATE TABLE performer (  
name VARCHAR **PRIMARY KEY**,  
type VARCHAR,  
country VARCHAR,  
style VARCHAR,  
founded INT,  
born INT,  
died INT  
);

# What is the syntax of the CREATE TABLE statement?

- Primary key declared in separate clause

CREATE TABLE performer (

name VARCHAR,

type VARCHAR,

country VARCHAR,

style VARCHAR,

founded INT,

born INT,

died INT,

**PRIMARY KEY (name)**

);

## How are primary key, partition key, and clustering columns defined?

- Simple partition key, no clustering columns  
PRIMARY KEY ( partition\_key\_column )
- Composite partition key, no clustering columns  
PRIMARY KEY ( ( partition\_key\_col1, ..., partition\_key\_colN ) )
- Simple partition key and clustering columns  
PRIMARY KEY ( partition\_key\_column,  
clustering\_column1, ..., clustering\_columnM )
- Composite partition key and clustering columns  
PRIMARY KEY ( ( partition\_key\_col1, ..., partition\_key\_colN ),  
clustering\_column1, ..., clustering\_columnM )

## Apache Cassandra

- Example: Can find all performers and albums for a given track title

```
CREATE TABLE albums_by_track (  
  track_title VARCHAR,  
  performer VARCHAR,  
  year INT,  
  album_title VARCHAR,  
  PRIMARY KEY  
  (track_title, performer, year, album_title)  
);
```



## Apache Cassandra

- Example: Can find a performer, genre, and all track numbers and titles for a given album title and year

```
CREATE TABLE tracks_by_album (  
  album_title VARCHAR,  
  year INT,  
  performer VARCHAR STATIC,  
  genre VARCHAR STATIC,  
  number INT,  
  track_title VARCHAR,  
  PRIMARY KEY  
  ((album_title, year), number)  
);
```



## What CQL data types are available?

CQL Type	Constants	Description
ASCII	strings	US-ASCII character string
BIGINT	integers	64-bit signed long
BLOB	blobs	Arbitrary bytes (no validation), expressed as hexadecimal
BOOLEAN	booleans	true or false
COUNTER	integers	Distributed counter value (64-bit long)
DECIMAL	integers, floats	Variable-precision decimal
DOUBLE	integers	64-bit IEEE-754 floating point
FLOAT	integers, floats	32-bit IEEE-754 floating point
INET	strings	IP address string in IPv4 or IPv6 format*
INT	integers	32-bit signed integer
LIST	n/a	A collection of one or more ordered elements
MAP	n/a	A JSON-style array of literals: { literal : literal, literal : literal ... }
SET	n/a	A collection of one or more elements
TEXT	strings	UTF-8 encoded string
TIMESTAMP	integers, strings	Date plus time, encoded as 8 bytes since epoch
TUPLE	n/a	Up to 32k fields
UUID	uuids	A UUID in standard UUID format
TIMEUUID	uuids	Type 1 UUID only (CQL 3)
VARCHAR	strings	UTF-8 encoded string
VARINT	integers	Arbitrary-precision integer



## What CQL data types are available?

CQL Type	Constants	Description
ASCII	strings	US-ASCII character string
BIGINT	integers	64-bit signed long
BLOB	blobs	Arbitrary bytes (no validation), expressed as hexadecimal
BOOLEAN	booleans	true or false
COUNTER	integers	Distributed counter value (64-bit long)
DECIMAL	integers, floats	Variable-precision decimal
DOUBLE	integers	64-bit IEEE-754 floating point
FLOAT	integers, floats	32-bit IEEE-754 floating point
INET	strings	IP address string in IPv4 or IPv6 format*
INT	integers	32-bit signed integer
LIST	n/a	A collection of one or more ordered elements
MAP	n/a	A JSON-style array of literals: { literal : literal, literal : literal ... }
SET	n/a	A collection of one or more elements
TEXT	strings	UTF-8 encoded string
TIMESTAMP	integers, strings	Date plus time, encoded as 8 bytes since epoch
TUPLE	n/a	Up to 32k fields
→ UUID	uuids	A UUID in standard UUID format
TIMEUUID	uuids	Type 1 UUID only (CQL 3)
VARCHAR	strings	UTF-8 encoded string
VARINT	integers	Arbitrary-precision integer

## Apache Cassandra – UUID y TIMEUUID

- Ids
  - UUID and TIMEUUID are universally unique identifiers
    - Generated programmatically
  - UUID
    - Format
      - hex{8}-hex{4}-hex{4}-hex{4}-hex{12}
      - 52b11d6d-16e2-4ee2-b2a9-5ef1e9589328
    - Used to assign conflict-free (unique) identifiers to data objects
    - Numeric range so vast that duplication is statistically all but imposible
    - CQL function uuid() generates a new UUID
  - UUID data type supports Version 4 UUIDs
    - Randomly generated sequence of 32 hex digits separated by dashes
    - 52b11d6d-16e2-4ee2-b2a9-5ef1e9589328



## What CQL data types are available?

CQL Type	Constants	Description
ASCII	strings	US-ASCII character string
BIGINT	integers	64-bit signed long
BLOB	blobs	Arbitrary bytes (no validation), expressed as hexadecimal
BOOLEAN	booleans	true or false
COUNTER	integers	Distributed counter value (64-bit long)
DECIMAL	integers, floats	Variable-precision decimal
DOUBLE	integers	64-bit IEEE-754 floating point
FLOAT	integers, floats	32-bit IEEE-754 floating point
INET	strings	IP address string in IPv4 or IPv6 format*
INT	integers	32-bit signed integer
LIST	n/a	A collection of one or more ordered elements
MAP	n/a	A JSON-style array of literals: { literal : literal, literal : literal ... }
SET	n/a	A collection of one or more elements
TEXT	strings	UTF-8 encoded string
TIMESTAMP	integers, strings	Date plus time, encoded as 8 bytes since epoch
TUPLE	n/a	Up to 32k fields
UUID	uuids	A UUID in standard UUID format
→ TIMEUUID	uuids	Type 1 UUID only (CQL 3)
VARCHAR	strings	UTF-8 encoded string
VARINT	integers	Arbitrary-precision integer

## Apache Cassandra – TIMEUUID

- Time Ids

- TIMEUUID data type supports Version 1 UUIDs
  - Embeds a time value within a UUID
  - Generated using time (60 bits), a clock sequence number (14 bits), and MAC address (48 bits)
    - 1be43390-9fe4-11e3-8d05-425861b86ab6
  - CQL function now() generates a new TIMEUUID
- Time can be extracted from TIMEUUID
  - CQL function `dateOf()` extracts the embedded timestamp as a date
- TIMEUUID values in clustering columns or in column names are ordered based on time
  - DESC order on TIMEUUID lists most recent data first

## Apache Cassandra – UUID y TIMEUUID

- Ids, Example
  - Users are identified by UUID
  - User activities (i.e., rating a track) are identified by TIMEUUID
    - A user may rate the same track multiple times
    - Activities are ordered by the time component of TIMEUUID

```
CREATE TABLE track_ratings_by_user (  
  user UUID,  
  activity TIMEUUID,  
  rating INT,  
  album_title VARCHAR,  
  album_year INT,  
  track_title VARCHAR,  
  PRIMARY KEY (user, activity)  
) WITH CLUSTERING ORDER BY (activity DESC);
```



## What CQL data types are available?

CQL Type	Constants	Description
ASCII	strings	US-ASCII character string
BIGINT	integers	64-bit signed long
BLOB	blobs	Arbitrary bytes (no validation), expressed as hexadecimal
BOOLEAN	booleans	true or false
COUNTER	integers	Distributed counter value (64-bit long)
DECIMAL	integers, floats	Variable-precision decimal
DOUBLE	integers	64-bit IEEE-754 floating point
FLOAT	integers, floats	32-bit IEEE-754 floating point
INET	strings	IP address string in IPv4 or IPv6 format*
INT	integers	32-bit signed integer
LIST	n/a	A collection of one or more ordered elements
MAP	n/a	A JSON-style array of literals: { literal : literal, literal : literal ... }
SET	n/a	A collection of one or more elements
TEXT	strings	UTF-8 encoded string
→ TIMESTAMP	integers, strings	Date plus time, encoded as 8 bytes since epoch
TUPLE	n/a	Up to 32k fields
UUID	uuids	A UUID in standard UUID format
TIMEUUID	uuids	Type 1 UUID only (CQL 3)
VARCHAR	strings	UTF-8 encoded string
VARINT	integers	Arbitrary-precision integer



## Apache Cassandra - TIMESTAMP

- **TIMESTAMP** holds date and time
  - 64-bit integer representing a number of milliseconds since January 1 1970 at 00:00:00 GMT
  - Entered as
    - 64-bit integer
    - String literal in the ISO 8601 format
      - 1979-12-18 08:12:51-0400
      - 2014-02-27
      - Other variations are allowed
  - Displayed in cqlsh as
    - yyyy-mm-dd HH:mm:ssZ



## What CQL data types are available?

CQL Type	Constants	Description
ASCII	strings	US-ASCII character string
BIGINT	integers	64-bit signed long
BLOB	blobs	Arbitrary bytes (no validation), expressed as hexadecimal
BOOLEAN	booleans	true or false
→ COUNTER	integers	Distributed counter value (64-bit long)
DECIMAL	integers, floats	Variable-precision decimal
DOUBLE	integers	64-bit IEEE-754 floating point
FLOAT	integers, floats	32-bit IEEE-754 floating point
INET	strings	IP address string in IPv4 or IPv6 format*
INT	integers	32-bit signed integer
LIST	n/a	A collection of one or more ordered elements
MAP	n/a	A JSON-style array of literals: { literal : literal, literal : literal ... }
SET	n/a	A collection of one or more elements
TEXT	strings	UTF-8 encoded string
TIMESTAMP	integers, strings	Date plus time, encoded as 8 bytes since epoch
TUPLE	n/a	Up to 32k fields
UUID	uuids	A UUID in standard UUID format
TIMEUUID	uuids	Type 1 UUID only (CQL 3)
VARCHAR	strings	UTF-8 encoded string
VARINT	integers	Arbitrary-precision integer

# Apache Cassandra - COUNTER

- Cassandra supports distributed counters
  - Useful for tracking a count
  - Counter column stores a number that can only be updated
    - Incremented or decremented
    - Cannot assign an initial value to a counter (initial value is 0)
  - Counter column **cannot be part of a primary key**
  - If a table has a counter column, **all non-counter columns must be part of a primary key**

## Apache Cassandra - COUNTER

- COUNTER example:

```
CREATE TABLE ratings_by_track (  
  album_title VARCHAR,  
  album_year INT,  
  track_title VARCHAR,  
  num_ratings COUNTER,  
  sum_ratings COUNTER,  
  PRIMARY KEY (album_title, album_year, track_title)  
);
```

# What are special properties of the COUNTER data type?

- Performance considerations
  - **Read** is as **efficient** as for **non-counter columns**
  - **Update** is fast but **slightly slower** than an update for **non-counter columns**
    - A read is required before a write can be performed
- Accuracy considerations
  - If a counter update is timed out, a client application cannot simply retry a “failed” counter update as the timed-out update may have been persisted
    - Counter update is not an idempotent operation
    - Running an increment twice is not the same as running it once



## What CQL data types are available?

CQL Type	Constants	Description
ASCII	strings	US-ASCII character string
BIGINT	integers	64-bit signed long
BLOB	blobs	Arbitrary bytes (no validation), expressed as hexadecimal
BOOLEAN	booleans	true or false
COUNTER	integers	Distributed counter value (64-bit long)
DECIMAL	integers, floats	Variable-precision decimal
DOUBLE	integers	64-bit IEEE-754 floating point
FLOAT	integers, floats	32-bit IEEE-754 floating point
INET	strings	IP address string in IPv4 or IPv6 format*
INT	integers	32-bit signed integer
→ LIST	n/a	A collection of one or more ordered elements
→ MAP	n/a	A JSON-style array of literals: { literal : literal, literal : literal ... }
→ SET	n/a	A collection of one or more elements
TEXT	strings	UTF-8 encoded string
TIMESTAMP	integers, strings	Date plus time, encoded as 8 bytes since epoch
TUPLE	n/a	Up to 32k fields
UUID	uuids	A UUID in standard UUID format
TIMEUUID	uuids	Type 1 UUID only (CQL 3)
VARCHAR	strings	UTF-8 encoded string
VARINT	integers	Arbitrary-precision integer

## Apache Cassandra – Collection Columns

- Collection columns are multi-valued columns
  - Designed to store discrete sets of data (e.g., tags for a blog post)
    - A collection is retrieved in its entirety
  - 64,000 - maximum number of elements in a collection
    - In practice – dozens or hundreds
  - 64 KB - maximum size of each collection element
    - In practice – much smaller
- Collection columns
  - cannot be part of a primary key
  - cannot be part of a partition key
  - cannot be used as a clustering column
  - cannot nest inside of another collection

## Apache Cassandra – Collection Columns

- Set – typed collection of unique values
  - keywords SET<VARCHAR>
    - Ordered by values
    - No duplicates
- List – typed collection of non-unique values
  - songwriters LIST<VARCHAR>
    - Ordered by position
    - Duplicates are allowed
- Map – typed collection of key-value pairs
  - tracks MAP<INT,VARCHAR>
    - Ordered by keys
    - Unique keys but not values



## Apache Cassandra – User-defined types

- User-defined types group related fields of information
  - Represents related data in a single table, instead of multiple, separate tables
  - Uses any data type, including collections and other user-defined types
  - Reserved words cannot be used as a name for a user-defined type
    - byte
    - smallint
    - complex
    - enum
    - date
    - interval
    - macaddr
    - bitstring

## Apache Cassandra - UDF

```
CREATE TYPE track (  
  album_title VARCHAR,  
  album_year INT,  
  track_title VARCHAR,  
);
```

## Apache Cassandra - UDF

- Table columns can be user-defined types
  - Requires the use of the **frozen keyword** in C\* 2.1
  - A user-defined type can be used as a data type for a collection

```
CREATE TABLE musicdb.track_ratings_by_user (
```

```
user UUID,
```

```
activity TIMEUUID,
```

```
rating INT,
```

```
song frozen <track>, -- To force the update of the full record
```

```
PRIMARY KEY (user, activity)
```

```
) WITH CLUSTERING ORDER BY (activity DESC);
```

## Apache Cassandra – ALTER TYPE

- ALTER TYPE can change a user-defined type
  - Change the type of a field
    - ALTER TYPE track ALTER album\_title TYPE BLOB;
    - Types must be compatible
  - Add a field to a type
    - ALTER TYPE track ADD track\_number INT;
  - Rename a field of a type
    - ALTER TYPE track RENAME album\_year TO year;
  - Rename a user-defined type
    - ALTER TYPE track RENAME TO song;

## Apache Cassandra – DROP TYPE

- DROP TYPE removes a user-defined type
  - Cannot drop a user-defined type that is in use by a table or another type

DROP TYPE track;



## What CQL data types are available?

CQL Type	Constants	Description
ASCII	strings	US-ASCII character string
BIGINT	integers	64-bit signed long
BLOB	blobs	Arbitrary bytes (no validation), expressed as hexadecimal
BOOLEAN	booleans	true or false
COUNTER	integers	Distributed counter value (64-bit long)
DECIMAL	integers, floats	Variable-precision decimal
DOUBLE	integers	64-bit IEEE-754 floating point
FLOAT	integers, floats	32-bit IEEE-754 floating point
INET	strings	IP address string in IPv4 or IPv6 format*
INT	integers	32-bit signed integer
LIST	n/a	A collection of one or more ordered elements
MAP	n/a	A JSON-style array of literals: { literal : literal, literal : literal ... }
SET	n/a	A collection of one or more elements
TEXT	strings	UTF-8 encoded string
TIMESTAMP	integers, strings	Date plus time, encoded as 8 bytes since epoch
TUPLE	n/a	Up to 32k fields
UUID	uuids	A UUID in standard UUID format
TIMEUUID	uuids	Type 1 UUID only (CQL 3)
VARCHAR	strings	UTF-8 encoded string
VARINT	integers	Arbitrary-precision integer

## Apache Cassandra - TUPLE

- Tuples hold fixed-length sets of typed positional fields
  - Convenient alternative to creating a user-defined type
  - Accommodates up to 32768 fields, but generally only use a few
  - **Useful when prototyping**
  - Must use the frozen keyword in C\* 2.1
  - Tuples **can be nested** in other tuples

# Apache Cassandra - TUPLE

```
CREATE TABLE user (  
id UUID PRIMARY KEY,  
email text,  
equalizer frozen<tuple<float,float,float,float,float,  
float,float,float,float,float>>,  
name text,  
preferences set<text>  
);
```



# Apache Cassandra – ORDER BY

- CLUSTERING ORDER BY defines how data values in clustering columns are ordered (ASC or DESC) in a table
  - ASC is the default order for all clustering columns
  - When retrieving data, the default order or the order specified by a CLUSTERING ORDER BY clause is used
- **The order can be reversed in a query using the ORDER BY clause**

## Apache Cassandra – ALTER TABLE

- ALTER TABLE manipulates the table metadata
  - Adding a column
    - ALTER TABLE album ADD cover\_image VARCHAR;
  - Changing a column data type
    - Types must be compatible
    - **Clustering and indexed columns are not supported**
      - ALTER TABLE album ALTER cover\_image TYPE BLOB;
  - Dropping a column
    - **PRIMARY KEY columns are not supported**
      - ALTER TABLE album DROP cover\_image;

# Apache Cassandra – DROP TABLE

- DROP TABLE removes a table (all data in the table is lost)
  - DROP TABLE album;

## What is a secondary index?

- Tables are indexed on columns in a primary key
  - Search on a partition key is very efficient
  - Search on a partition key and clustering columns is very efficient
  - **Search on other columns is not supported**
    - **Last version of Cassandra relax this constraint**
- Secondary indexes
  - Can index additional columns to enable searching by those columns
    - one column per index
  - Cannot be created for
    - counter columns
    - static columns

## What is a secondary index?

```
CREATE TABLE performer (  
  name VARCHAR,  
  type VARCHAR,  
  country VARCHAR,  
  style VARCHAR,  
  founded INT,  
  born INT,  
  died INT,  
  PRIMARY KEY (name)  
);  
CREATE INDEX performer_style_key ON performer (style);
```

## What is a secondary index?

- `DROP INDEX performer_style_key;`

## When do you want to use a secondary index?

- Secondary indexes are for searching convenience
  - Use with low-cardinality columns
    - **Columns that may contain a relatively small set of distinct values**
      - For example, there are many artists but only a few dozen music styles
      - Allows searching for all artists for a specified style (a potentially expensive query because it may return a large result set)
  - Use with smaller datasets or when prototyping
- Do not use
  - On high-cardinality columns
  - On counter column tables
  - On a frequently updated or deleted columns
  - To look for a row in a large partition unless narrowly queried
    - e.g., search on both a partition key and an indexed column

Time for Exercises