

Module 4

NoSQL Databases

Learning Objectives

- Understand the Cassandra data model
- Introduce cqlsh
- Understand and use the DDL subset of CQL
- Introduce DevCenter
- **Understand and use the DML subset of CQL**
- ~~Understand basics of data modeling (Challenge)~~

What is the syntax of the INSERT statement?

```
INSERT INTO table_name (column1, column2 ...)
```

```
VALUES (value1, value2 ...)
```

- Inserts a row into a table
 - Must specify columns to insert values into
 - Primary key columns are mandatory (identify the row)
 - Other columns do not have to have values
 - Non-existent 'values' do not take up space
- Atomicity and isolation
 - Inserts are atomic
 - All values of a row are inserted or none
 - Inserts are isolated
 - Two inserts with the same values in primary key columns will not interfere-executed one after another

What is the syntax of the INSERT statement?

- To insert a row into a table

```
CREATE TABLE albums_by_performer (  
  performer VARCHAR,  
  year INT,  
  title VARCHAR,  
  genre VARCHAR,  
  PRIMARY KEY (performer, year, title)  
) WITH CLUSTERING ORDER BY (year DESC, title ASC);
```

```
INSERT INTO albums_by_performer (performer,year,title,genre)  
VALUES ('The Beatles', 1966, 'Revolver', 'Rock');  
INSERT INTO albums_by_performer (performer, year, title)  
VALUES ('The Beatles', 1995, 'Beatlemania');
```

What is the syntax of the INSERT statement?

performer	year	title	Genre
The Beatles	1995	Beatlemania	
The Beatles	1966	Revolver	Rock

Apache Cassandra - UPDATE

```
UPDATE <keyspace>.<table>
```

```
SET column_name1 = value, column_name2 = value,
```

```
WHERE primary_key_column = value
```

- Updates columns in an existing row
 - Row must be identified by values in primary key columns
 - Primary key columns cannot be updated
 - An existing value is replaced with a new value
 - A new value is added if a value for a column did not exist before
- Atomicity and isolation
 - Updates are atomic
 - All values of a row are updated or none
 - Updates are isolated
 - Two updates with the same values in primary key columns will not interfere – executed one after another

Apache Cassandra - UPDATE

- To update a row in a table

```
UPDATE albums_by_performer
```

```
SET genre = 'Rock'
```

```
WHERE performer = 'The Beatles' AND
```

```
year = 1995 AND
```

```
title = 'Beatlemania';
```

Apache Cassandra - UPDATE

- Before update

performer	year	title	Genre
The Beatles	1995	Beatlemania	
The Beatles	1966	Revolver	Rock

- After update

performer	year	title	Genre
The Beatles	1995	Beatlemania	Rock
The Beatles	1966	Revolver	Rock

Apache Cassandra - UPDATE

- UPdate + inSERT
 - Both UPDATE and INSERT are write operations
 - No reading before writing
- Term “upsert” denotes the following behavior
 - INSERT updates or overwrites an existing row
 - When inserting a row in a table that already has another row with the same values in primary key columns
 - UPDATE inserts a new row
 - When a to-be-updated row, identified by values in primary key columns, does not exist
 - Upserts are legal and do not result in error or warning messages

Apache Cassandra - UPDATE

- Introduces a new clause IF NOT EXISTS for inserts
 - Insert operation executes if a row with the same primary key does not exist
 - Uses a consensus algorithm called Paxos to ensure inserts are done serially
 - Multiple messages are passed between coordinator and replicas with a **large performance penalty**
 - [applied] column returns true if row does not exist and insert executes
 - [applied] column is false if row exists and the existing row will be returned

Apache Cassandra – IF NOT EXISTS

```
INSERT INTO albums_by_performer (performer,year,title)  
VALUES ('The Beatles', 1966, 'Revolver') IF NOT EXISTS;
```

[applied]

true

```
INSERT INTO albums_by_performer (performer,year,title)  
VALUES ('The Beatles', 1966, 'Revolver') IF NOT EXISTS;
```

[applied]

performer

year

false

The Beatles

1966

Apache Cassandra - IF

- Update uses IF to verify the value for column(s) before execution
 - [applied] column returns true if condition(s) matches and update written
 - [applied] column is false if condition(s) do not match and the current row will be returned

Apache Cassandra

UPDATE albums_by_performer SET year = 1968 WHERE performer = 'The Beatles'
IF title = 'Revolver';

[applied]

true

UPDATE albums_by_performer SET year = 1968 WHERE performer = 'The Beatles'
IF title = 'Revolver' AND year = 1967;

[applied]

performer

year

false

The Beatles

1966

Apache Cassandra - TTL

- Time-to-live (TTL) defines expiring columns
- INSERT and UPDATE can optionally assign data values a time-to-live
 - TTL is specified in seconds
 - Expired columns/values are eventually deleted
 - With no TTL specified, columns/values never expire
- TTL is useful for automatic deletion
 - When data gets outdated after some time
 - When only most recent data is needed
 - Older data may be archived elsewhere by a background process
 - Helps keep the size of a table and its partitions manageable
 - Restricts the data view to most recent data

Apache Cassandra – TTL

- To store a row for 86400 seconds (1 day)
 - Re-inserting the same row before it expires will overwrite TTL

```
INSERT INTO track_ratings_by_user  
(user,activity,rating,album_title,album_year,track_title)  
VALUES (uuid(),now(),5,'Revolver',1966,'Yellow  
Submarine')  
USING TTL 86400;
```

Apache Cassandra

- To store a column value for 30 seconds
 - Only column 'rating' for this row is affected by TTL

```
UPDATE track_ratings_by_user
```

```
USING TTL 30
```

```
SET rating = 0
```

```
WHERE user = 52b11d6d-16e2-4ee2-b2a9-5ef1e9589328 AND
```

```
activity = dbf3fbfc-9fe4-11e3-8d05-425861b86ab6;
```


Apache Cassandra - DELETE

- Deletes a partition, a row or specified columns in a row
 - Row must be identified by values in primary key columns
 - Primary key columns cannot be deleted without deleting the whole row
- To delete a partition from a table
- To delete a row from a table
- To delete a column from a table row

Apache Cassandra - DELETE

- Deletes a partition, a row or specified columns in a row
 - Row must be identified by values in primary key columns
 - Primary key columns cannot be deleted without deleting the whole row
- **To delete a partition from a table**

```
DELETE FROM track_ratings_by_user
```

```
WHERE user = 52b11d6d-16e2-4ee2-b2a9-5ef1e9589328;
```

Apache Cassandra - DELETE

- Deletes a partition, a row or specified columns in a row
 - Row must be identified by values in primary key columns
 - Primary key columns cannot be deleted without deleting the whole row
- To delete a partition from a table
- **To delete a row from a table**

```
DELETE FROM track_ratings_by_user
```

```
WHERE user = 52b11d6d-16e2-4ee2-b2a9-5ef1e9589328 AND activity = dbf3fbfc-9fe4-11e3-8d05-425861b86ab6;
```

Apache Cassandra - DELETE

- Deletes a partition, a row or specified columns in a row
 - Row must be identified by values in primary key columns
 - Primary key columns cannot be deleted without deleting the whole row
- To delete a partition from a table
- To delete a row from a table
- **To delete a column from a table row**

DELETE rating FROM track_ratings_by_user

WHERE user = 52b11d6d-16e2-4ee2-b2a9-5ef1e9589328 AND activity = dbf3fbfc-9fe4-11e3-8d05-425861b86ab6;

Apache Cassandra - TRUNCATE

- TRUNCATE removes all rows in a table
 - The table definition (schema) is not affected

```
TRUNCATE track_ratings_by_user;
```

Apache Cassandra - LIST

- CQL list – defining and inserting
 - Collection column cannot be part of a primary key

```
CREATE TABLE song (  
id UUID PRIMARY KEY,  
title VARCHAR,  
songwriters LIST<VARCHAR>  
);  
  
INSERT INTO song (id, title, songwriters)  
VALUES (uuid(),  
'I Want to Hold Your Hand', ['John', 'Paul']);
```

Apache Cassandra - LIST

- CQL list – defining and inserting
 - Collection column cannot be part of a primary key

```
CREATE TABLE song (  
id UUID PRIMARY KEY,  
title VARCHAR,  
songwriters LIST<VARCHAR>  
);  
INSERT INTO song (id, title, songwriters)  
VALUES (uuid(),  
'I Want to
```

id	songwriters	title
252608cb-0f56-4cf3-82ee-b7fe00f3920f	['John', 'Paul']	I Want to Hold Your Hand

Apache Cassandra - LIST

id	songwriters	title
252608cb-0f56-4cf3-82ee-b7fe00f3920f	['John', 'Paul']	I Want to Hold Your Hand

UPDATE song SET songwriters = songwriters +
['Paul', 'Jonathan']

WHERE id = 252608cb-0f56-4cf3-82ee-b7fe00f3920f;

id	songwriters	title
252608cb-0f56-4cf3-82ee-b7fe00f3920f	['John', 'Paul', 'Paul', 'Jonathan']	I Want to Hold Your Hand

Apache Cassandra - LIST

id	songwriters	title
252608cb-0f56-4cf3-82ee-b7fe00f3920f	['John', 'Paul', 'Paul', 'Jonathan']	I Want to Hold Your Hand

```
UPDATE song SET songwriters = ['Patrick'] + songwriters  
WHERE id = 252608cb-0f56-4cf3-82ee-b7fe00f3920f;
```

id	songwriters	title
252608cb-0f56-4cf3-82ee-b7fe00f3920f	['Patrick', 'John', 'Paul', 'Paul', 'Jonathan']	I Want to Hold Your Hand

Apache Cassandra - LIST

id	songwriters	title
252608cb-0f56-4cf3-82ee-b7fe00f3920f	['Patrick', 'John', 'Paul', 'Paul', 'Jonathan']	I Want to Hold Your Hand

```
UPDATE song SET songwriters[3] = 'Ringo'
```

```
WHERE id = 252608cb-0f56-4cf3-82ee-b7fe00f3920f;
```

id	songwriters	title
252608cb-0f56-4cf3-82ee-b7fe00f3920f	['Patrick', 'John', 'Paul', 'Ringo', 'Jonathan']	I Want to Hold Your Hand

Apache Cassandra - LIST

id	songwriters	title
252608cb-0f56-4cf3-82ee-b7fe00f3920f	['Patrick', 'John', 'Paul', 'Ringo', 'Jonathan']	I Want to Hold Your Hand

UPDATE song SET songwriters = songwriters -

['Patrick', 'Jonathan', 'Ringo']

WHERE id = 252608cb-0f56-4cf3-82ee-b7fe00f3920f;

id	songwriters	title
252608cb-0f56- ...	['John', 'Paul']	I Want to Hold Your Hand

Apache Cassandra - LIST

id	songwriters	title
252608cb-0f56- ...	['John', 'Paul']	I Want to Hold Your Hand

DELETE songwriters[0], songwriters[1] FROM song
WHERE id = 252608cb-0f56-4cf3-82ee-b7fe00f3920f;

id	songwriters	title
252608cb-0f56- ...		I Want to Hold Your Hand

Apache Cassandra - MAP

- CQL map – defining and inserting
 - Collection column cannot be part of a primary key

```
CREATE TABLE album (  
  title VARCHAR,  
  year INT,  
  tracks MAP<INT,VARCHAR>,  
  PRIMARY KEY ((title, year))  
);  
  
INSERT INTO album (title, year, tracks)  
VALUES ('Revolver', 1966, {1: 'Taxman', 2: 'Eleanor Rigby'});
```

title	year	tracks
Revolver	1966	{1: 'Taxman', 2: 'Eleanor Rigby'}

Apache Cassandra - MAP

title	year	tracks
Revolver	1966	{1: 'Taxman', 2: 'Eleanor Rigby'}

```
UPDATE album SET tracks[14] = 'Yellow Submarine'  
WHERE title = 'Revolver' AND year = 1966;
```

title	year	tracks
Revolver	1966	{1: 'Taxman', 2: 'Eleanor Rigby', 14: 'Yellow Submarine'}

Apache Cassandra - MAP

title	year	tracks
Revolver	1966	{1: 'Taxman', 2: 'Eleanor Rigby', 14: 'Yellow Submarine'}

```
UPDATE album SET tracks[14] = 'Tomorrow Never Knows'  
WHERE title = 'Revolver' AND year = 1966;
```

title	year	tracks
Revolver	1966	{1: 'Taxman', 2: 'Eleanor Rigby', 14: 'Tomorrow Never Knows'}

Apache Cassandra - MAP

title	year	tracks
Revolver	1966	{1: 'Taxman', 2: 'Eleanor Rigby', 14: 'Tomorrow Never Knows'}

DELETE tracks[14] FROM album
WHERE title = 'Revolver' AND year = 1966;

title	year	tracks
Revolver	1966	{1: 'Taxman', 2: 'Eleanor Rigby'}

Apache Cassandra - MAP

title	year	tracks
Revolver	1966	{1: 'Taxman', 2: 'Eleanor Rigby'}

DELETE tracks FROM album

WHERE title = 'Revolver' AND year = 1966;

title	year	tracks
Revolver	1966	

Apache Cassandra - TUPLE

```
CREATE TABLE user (  
id UUID PRIMARY KEY,  
email text,  
name text,  
preferences set<text>,  
equalizer frozen<tuple<float,float,float,float,float,  
float,float,float,float,float>>  
>;  
INSERT INTO user (id, equalizer)  
VALUES (uuid(),(3.0, 6.0, 9.0, 7.0, 6.0, 5.0, 7.0, 9.0, 11.0, 8.0));
```

Apache Cassandra - TUPLE

```
CREATE TABLE user (  
id UUID PRIMARY KEY,  
email text,  
name text,  
preferences set<text>,  
equalizer frozen<tuple<float,float,float,float,float,  
float,float,float,float,float>>  
);
```

```
INSERT IN  
VALUES (u
```

id	equalizer
62d4f220-5361-...	(3.0, 6.0, 9.0, 7.0, 6.0, 5.0, 7.0, 9.0, 11.0, 8.0)

Apache Cassandra - TUPLE

id	equalizer
62d4f220-5361-...	(3.0, 6.0, 9.0, 7.0, 6.0, 5.0, 7.0, 9.0, 11.0, 8.0)

UPDATE user SET equalizer =

(4.0, 1.6, -1.8, -5.6, -0.7, 0.9, 2.9, 4.3, 4.3, 4.3)

WHERE id = 6ed4f220-5361-11e4-8d89-c971d060d947;

id	equalizer
62d4f220-5361-...	(4.0, 1.6, -1.8, -5.6, -0.7, 0.9, 2.9, 4.3, 4.3, 4.3)

Apache Cassandra - TUPLE

id	equalizer
62d4f220-5361-...	(4.0, 1.6, -1.8, -5.6, -0.7, 0.9, 2.9, 4.3, 4.3, 4.3)

DELETE equalizer from user

WHERE id = 6ed4f220-5361-11e4-8d89-c971d060d947

id	equalizer
62d4f220-5361-...	

Apache Cassandra - BATCH

- BATCH statement combines multiple INSERT, UPDATE, and DELETE statements into a single logical operation
 - Saves on client-server and coordinator-replica communication
 - Atomic operation
 - If any statement in the batch succeeds, all will
 - No batch isolation
 - Other “transactions” can read and write data being affected by partially executed batch

Apache Cassandra - BATCH

BEGIN BATCH

```
DELETE FROM albums_by_performer
```

```
WHERE performer = 'The Beatles' AND year = 1966 AND title = 'Revolver';
```

```
INSERT INTO albums_by_performer (performer, year, title, genre)
```

```
VALUES ('The Beatles', 1966, 'Revolver', 'Rock');
```

APPLY BATCH;

Apache Cassandra - LIGHTWEIGHT

- Lightweight transactions in batch
 - Batch will execute only if conditions for all lightweight transactions are met
 - All operations in batch will execute serially with the increased performance overhead

Apache Cassandra - LIGHTWEIGHT

BEGIN BATCH

```
UPDATE user SET lock = true IF lock = false;
```

```
WHERE performer = 'The Beatles' AND year = 1966 AND title = 'Revolver';
```

```
INSERT INTO albums_by_performer (performer, year, title, genre)
```

```
VALUES ('The Beatles', 1966, 'Revolver', 'Rock');
```

```
UPDATE user SET lock = false;
```

APPLY BATCH;

Apache Cassandra - SELECT

- Retrieves rows from a table that satisfy an optional condition
 - SELECT – Which columns to retrieve?
 - FROM – Which table to retrieve from?
 - WHERE – What condition must rows satisfy?
 - ORDER BY – How to sort a result set?
 - LIMIT – How many rows to return?
 - ALLOW FILTERING – Is scanning over all partitions allowed?

Apache Cassandra - SELECT

SELECT select_expression

FROM keyspace_name.table_name

WHERE relation AND relation ...

ORDER BY (clustering_column (ASC | DESC)...)

LIMIT n

ALLOW FILTERING

Apache Cassandra - SELECT

- To retrieve all rows
 - `SELECT * FROM album;`
- To retrieve specific columns of all rows
 - `SELECT performer, title, year FROM album;`
- To retrieve a specific field from a user-defined type column
 - `SELECT performer.lastname FROM album;`
- To compute the number of rows in a table
 - `SELECT COUNT(*) FROM album;`

Apache Cassandra - SELECT

- Equality search – one partition
 - To retrieve one partition, values for all partition key columns must be specified
 - In a single-row partition, row = partition

```
CREATE TABLE tracks_by_album ( ...  
PRIMARY KEY ((album_title, year), number));  
  
SELECT album_title, year, number, track_title  
FROM tracks_by_album  
WHERE album_title = 'Revolver' AND year = 1966;
```

Apache Cassandra - SELECT

- Equality search – one partition

- To retrieve one partition. values for all partition key columns must be specified

• In a si

album_title	year	number	track_title
Revolver	1966	1	Taxman
Revolver	1966	2	Eleanor Rigby
...
Revolver	1966	14	Tomorrow Never Knows

SELECT album_title, year, number, track_title

FROM tracks_by_album

WHERE album_title = 'Revolver' AND year = 1966;

Apache Cassandra - SELECT

- Equality search – one row
 - To retrieve one row, values for all primary key columns must be specified
 - In a single-row partition, primary key = partition key

```
CREATE TABLE tracks_by_album ( ...  
PRIMARY KEY ((album_title, year), number));
```

```
SELECT album_title, year, number, track_title  
FROM tracks_by_album  
WHERE album_title = 'Revolver' AND year = 1966 AND  
number = 6;
```

Apache Cassandra - SELECT

- Equality search – one row
 - To retrieve one row, values for all primary key columns must be specified
 - In a single-row partition, primary key = partition key

```
CREATE TABLE tracks_by_album ( ...
```

```
PRIMARY
```

album_title	year	number	track_title
Revolver	1966	6	Yellow Submarine

```
FROM tracks_by_album
```

```
WHERE album_title = 'Revolver' AND year = 1966 AND  
number = 6;
```


Apache Cassandra - SELECT

- Equality search – subset of rows
 - To retrieve a subset of rows in a partition, values for all partition key columns and all clustering columns must be specified with the last clustering column value being a set
 - IN is only allowed on the **last clustering column** of a primary key

```
CREATE TABLE tracks_by_album ( ...  
PRIMARY KEY ((album_title, year), number));
```

```
SELECT album_title, year, number, track_title  
FROM tracks_by_album  
WHERE album_title = 'Revolver' AND year = 1966 AND  
number IN (2,6,7,14);
```

Apache Cassandra - SELECT

- Equality search – subset of rows

- To retrieve a subset of rows in a partition, values for all partition key columns and all clustered columns. This is done using a set

- IN is used

CREATE

PRIMARY

SELECT album_title, year, number, track_title

FROM tracks_by_album

WHERE album_title = 'Revolver' AND year = 1966 AND

number IN (2,6,7,14);

album_title	year	number	track_title
Revolver	1966	2	Eleanor Rigby
Revolver	1966	6	Yellow Submarine
Revolver	1966	7	She Said She Said
Revolver	1966	14	Tomorrow Never Knows

Apache Cassandra - SELECT

- Equality search – subset of rows
 - To retrieve a subset of rows in a partition, values for all partition key columns and one or more but not all clustering columns must be specified
 - Clustering columns in a predicate must constitute a **prefix of clustering columns** specified in the primary key definition

```
CREATE TABLE albums_by_performer ( ...  
PRIMARY KEY (performer, year, title));
```

```
SELECT title, year  
FROM albums_by_performer  
WHERE performer = 'The Beatles' AND year = 1970;
```

Apache Cassandra - SELECT

- Equality search – subset of rows

- To retrieve a subset of rows in a partition, values for all partition key columns and one or more but not all clustering columns must be specified
- Clustering columns in a predicate must constitute a **prefix of clustering columns** specified

	title	year
CREATE	At The Hollywood Bowl	1970
PRIMARY	Let It Be	1970
	The Beatles Christmas Album	1970

```
SELECT title, year  
FROM albums_by_performer  
WHERE performer = 'The Beatles' AND year = 1970;
```

Apache Cassandra - SELECT

- Equality search – multiple partitions
 - To retrieve multiple partitions, a set of values for a partition key must be specified using IN
 - IN is only allowed on the last column of a partition key

```
CREATE TABLE albums_by_performer ( ...  
PRIMARY KEY (performer, year, title));  
SELECT performer, title, year  
FROM albums_by_performer  
WHERE performer IN ('The Beatles', 'Deep Purple');
```

Apache Cassandra - SELECT

- Equality search – multiple partitions

- To retrieve data using IN

- IN is used

CREATE

PRIMARY

SELECT

FROM albums_by_performer

WHERE performer IN ('The Beatles', 'Deep Purple');

performer	title	year
The Beatles	Let It Be...Naked	2003
...
The Beatles	With The Beatles	1963
Deep Purple	Abandon	1998
...

Apache Cassandra - SELECT

- Range search
 - $>$, \geq , $<$, \leq
 - Can only a range search on a partition key using the token() function
 - Results are not meaningful for RandomPartitioner and Murmur3Partitioner
 - Allowed on only one clustering column in a predicate
 - This column should be defined later in the PRIMARY KEY clause than any other clustering column used in a predicate

WHERE token(key) \geq token(?) AND token(key) $<$ token(?)

Apache Cassandra - SELECT

```
CREATE TABLE tracks_by_album ( ...
```

```
PRIMARY KEY ((album_title, year), number));
```

```
SELECT album_title, year, number, track_title
```

```
FROM tracks_by_album
```

```
WHERE album_title = 'Revolver' AND year = 1966 AND number >= 6 AND number  
< 8;
```


Apache Cassandra - SELECT

```
CREATE TABLE tracks_by_album ( ...  
PRIMARY KEY ((album_title, year), number));
```

```
SELECT album_title, year, number, track_title  
FROM tracks_by_album
```

```
WHERE album_title = 'Revolver' AND year = 1966 AND number >= 6 AND number  
< 8;
```

album_title	year	number	track_title
Revolver	1966	6	Yellow Submarine
Revolver	1966	7	She Said She Said

Apache Cassandra – ALLOW FILTERING

- Allows scanning over all partitions
 - Predicate does not specify values for partition key columns
 - Relaxes the requirement that a partition key must be specified
 - Potentially expensive queries that may return large results
 - Use with caution
 - LIMIT clause is recommended
 - Predicate can have equality or inequality relations on clustering columns
- Return 7th tracks for the first 10 albums in the table
- Return the number of albums with 30 or more tracks

Apache Cassandra

- Allows scanning over all partitions
 - Predicate does not specify values for partition key columns
 - Relaxes the requirement that a partition key must be specified
 - Potentially expensive queries that may return large results
 - Use with caution
 - LIMIT clause is recommended
 - Predicate can have equality or inequality relations on clustering columns
- **Return 7th tracks for the first 10 albums in the table**

```
SELECT * FROM tracks_by_album  
WHERE number = 7 LIMIT 10 ALLOW FILTERING;
```
- Return the number of albums with 30 or more tracks

Apache Cassandra

- Allows scanning over all partitions
 - Predicate does not specify values for partition key columns
 - Relaxes the requirement that a partition key must be specified
 - Potentially expensive queries that may return large results
 - Use with caution
 - LIMIT clause is recommended
 - Predicate can have equality or inequality relations on clustering columns
- Return 7th tracks for the first 10 albums in the table
- **Return the number of albums with 30 or more tracks**

```
SELECT COUNT(*) FROM tracks_by_album  
WHERE number = 30 LIMIT 100000 ALLOW FILTERING;
```

Apache Cassandra – Indexes in Queries

- A predicate may involve only an indexed column
- A predicate may involve primary key and indexed columns
 - Useful to narrow a search in a large multi-row partition
- A predicate may involve multiple indexed columns
 - ALLOW FILTERING must be used

Apache Cassandra – Indexes in Queries

- **A predicate may involve only an indexed column**

```
CREATE INDEX performer_country_key ON performer (country);  
SELECT name FROM performer WHERE country = 'Iceland';
```

- A predicate may involve primary key and indexed columns
 - Useful to narrow a search in a large multi-row partition
- A predicate may involve multiple indexed columns
 - ALLOW FILTERING must be used

Apache Cassandra – Indexes in Queries

- A predicate may involve only an indexed column
- A predicate may involve primary key and indexed columns
 - Useful to narrow a search in a large multi-row partition
- **A predicate may involve multiple indexed columns**
 - ALLOW FILTERING must be used

```
CREATE INDEX performer_country_key ON performer (country);
```

```
CREATE INDEX performer_style_key ON performer (style);
```

```
SELECT name FROM performer
```

```
WHERE country = 'Iceland' AND style = 'Rock' ALLOW FILTERING;
```

Apache Cassandra – SELECT in Collections

- Searches on indexed collections uses the CONTAINS keyword
- Set, List, Map – Search for a value
- Map – Search for a key

Apache Cassandra

- Searches on indexed collections uses the CONTAINS keyword
- **Set, List, Map – Search for a value**

```
CREATE INDEX ON user (preferences);
```

```
SELECT id FROM user
```

```
WHERE preferences CONTAINS 'Rock';
```

- Map – Search for a key

Apache Cassandra – SELECT in Collections

- Searches on indexed collections uses the CONTAINS keyword
- Set, List, Map – Search for a value
- **Map – Search for a key**

```
CREATE INDEX ON album (tracks);
```

```
SELECT title, tracks FROM album  
WHERE tracks CONTAINS KEY 20;
```

Apache Cassandra – SELECT in Collections

- The column is treated as a blob and must search on all fields
 - User-defined type – Search all fields
 - Tuple – Search all fields

Apache Cassandra

- The column is treated as a blob and must search on all fields

- **User-defined type – Search all fields**

```
CREATE INDEX ON track_ratings_by_user (song);
```

```
SELECT * FROM track_ratings_by_user  
WHERE song = {album_title: 'Beatles For Sale',  
album_year: 1964,  
track_title: 'Cant Buy Me Love'};
```

- Tuple – Search all fields

Apache Cassandra – SELECT in Collections

- The column is treated as a blob and must search on all fields
 - User-defined type – Search all fields
 - **Tuple – Search all fields**

```
CREATE INDEX ON user (equalizer);
```

```
SELECT * FROM user
```

```
WHERE equalizer = (1.0, 2.0, 3.0, 4.0, 5.0,  
6.0, 7.0, 8.0, 9.0, 10.0);
```

Apache Cassandra – ORDER BY

- ORDER BY specifies how query results must be sorted
 - Allowed only on clustering columns
 - Default order is ASC or as defined by WITH CLUSTERING ORDER
 - Default order can be reversed for all clustering columns at once

Apache Cassandra – ORDER BY

```
CREATE TABLE tracks_by_album ( ...  
PRIMARY KEY ((album_title, year), number));  
  
SELECT album_title, year, number, track_title  
FROM tracks_by_album  
WHERE album_title = 'Revolver' AND year = 1966  
ORDER BY number DESC;
```

Apache

CREATE TABLE

PRIMARY

album_title	year	number	track_title
Revolver	1966	14	Tomorrow Never Knows
Revolver	1966	13	Got to Get You Into My Life
...
Revolver	1966	1	Taxman

SELECT album_title, year, number, track_title

FROM tracks_by_album

WHERE album_title = 'Revolver' AND year = 1966

ORDER BY number DESC;

Apache Cassandra - Functions

- TIMEUUID functions
 - `dateOf()` – extracts the timestamp as a date of a timeuuid column
 - `now()` – generates a new unique timeuuid
 - `minTimeuuid()` and `maxTimeuuid()` – return a UUID-like result given a conditional time component as an argument
 - `unixTimestampOf()` – extracts the “raw” timestamp of a timeuuid column as a 64-bit integer

Apache Cassandra - Functions

- TIMEUUID functions
- **dateOf()** – **extracts the timestamp as a date of a timeuuid column**
 - `SELECT dateOf(timeuuid_column), ... FROM ...;`
- `now()` – generates a new unique timeuuid
- `minTimeuuid()` and `maxTimeuuid()` – return a UUID-like result given a conditional time component as an argument
- `unixTimestampOf()` – extracts the “raw” timestamp of a timeuuid column as a 64-bit integer

Apache Cassandra - Functions

- TIMEUUID functions
- `dateOf()` – extracts the timestamp as a date of a timeuuid column
- **`now()` – generates a new unique timeuuid**
 - `INSERT INTO ... (timeuuid_column, ...) VALUES (now(), ...);`
- `minTimeuuid()` and `maxTimeuuid()` – return a UUID-like result given a conditional time component as an argument
- `unixTimestampOf()` – extracts the “raw” timestamp of a timeuuid column as a 64-bit integer

Apache Cassandra - Functions

- TIMEUUID functions
- `dateOf()` – extracts the timestamp as a date of a timeuuid column
- `now()` – generates a new unique timeuuid
- **`minTimeuuid()` and `maxTimeuuid()` – return a UUID-like result given a conditional time component as an argument**
 - `SELECT * FROM ... WHERE ... AND timeuuid_column > maxTimeuuid('2014-01-01 00:00+0000') AND timeuuid_column < minTimeuuid('2014-03-01 00:00+0000');`
- `unixTimestampOf()` – extracts the “raw” timestamp of a timeuuid column as a 64-bit integer

Apache Cassandra - Functions

- TIMEUUID functions
- `dateOf()` – extracts the timestamp as a date of a timeuuid column
- `now()` – generates a new unique timeuuid
- `minTimeuuid()` and `maxTimeuuid()` – return a UUID-like result given a conditional time component as an argument
- **`unixTimestampOf()` – extracts the “raw” timestamp of a timeuuid column as a 64-bit integer**
 - `SELECT unixTimestampOf(timeuuid_column), ... FROM ...;`

Apache Cassandra - BLOB

- Blob conversion functions
 - Series of `typeAsBlob()` and `blobAsType()` functions
 - The Cassandra blob data type represents a constant hexadecimal number defined as `0[xX](hex)+` where hex is a hexadecimal character, such as `[0-9a-fA-F]`.
 - For example, `0xcafe`.
 - The maximum theoretical size for a blob is 2 GB. The practical limit on blob size, however, is less than 1 MB. A blob type is suitable for storing a small image or short string.
- Token access function
 - `token()` function

Apache Cassandra - BLOB

- Blob conversion functions

- **Series of typeAsBlob() and blobAsType() functions**

```
SELECT varcharAsBlob(varchar_column), ... FROM ...;
```

```
SELECT blobAsBigint(blob_column), ... FROM ...;
```

```
CREATE TABLE bios ( user_name varchar PRIMARY KEY,  
    bio blob  
    );
```

```
INSERT INTO bios (user_name, bio) VALUES ('fred', bigintAsBlob(3));
```

- Token access function

- token() function

Apache Cassandra - BLOB

- Blob conversion functions
 - Series of `typeAsBlob()` and `blobAsType()` functions
- Token access function
 - **`token()` function**

```
SELECT * FROM ... WHERE token(partition_key) > token(2014);
```